

Design and Analysis
of Algorithms I

Graph Primitives

Dijkstra's Algorithm: The Basics

Single-Source Shortest Paths

Input: directed graph $G=(V, E)$. ($m=|E|$, $n=|V|$)

- each edge has non negative length l_e
- source vertex s

Output: for each $v \in V$, compute

$L(v) :=$ length of a shortest $s-v$ path in G

Assumption:

1. [for convenience] $\forall v \in V, \exists s \Rightarrow v$ path
2. [important] $l_e \geq 0 \quad \forall e \in E$

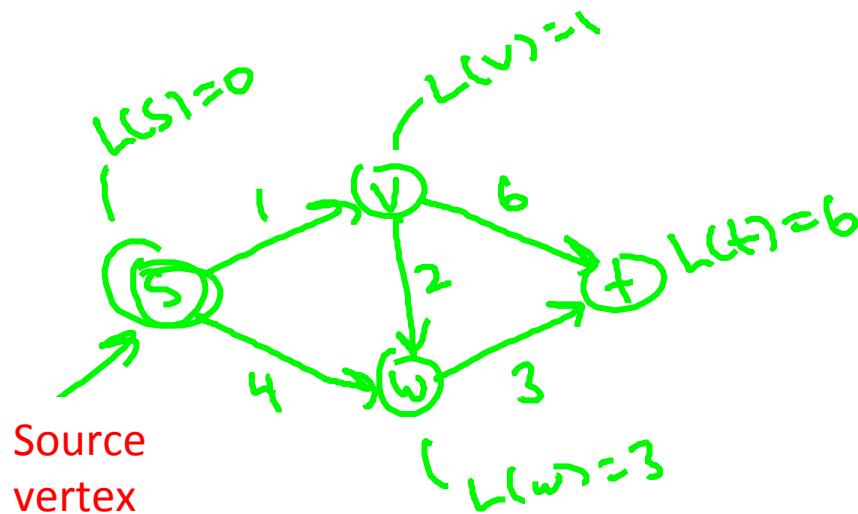
Length of path
= sum of edge lengths



Path length = 6

One of the following is the list of shortest-path distances for the nodes s, v, w, t , respectively. Which is it?

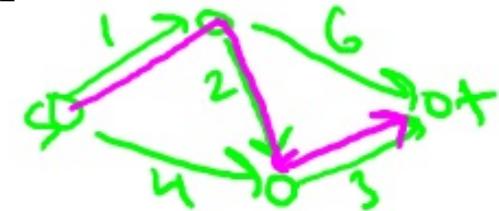
- 0,1,2,3
- 0,1,4,7
- 0,1,4,6
- 0,1,3,6



Why Another Shortest-Path Algorithm?

Question: doesn't BFS already compute shortest paths in linear time?

Answer: yes, IF $l_e = 1$ for every edge e .



Question: why not just replace each edge e by directed path of l_e unit length edges:



Answer: blows up graph too much

Solution: Dijkstra's shortest path algorithm.

This array
only to help
explanation!

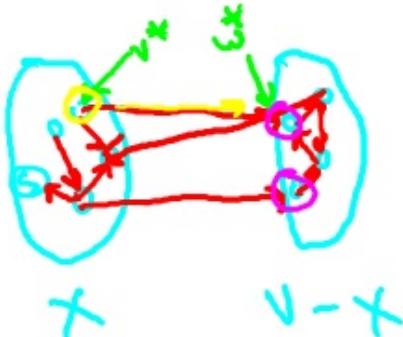
Dijkstra's Algorithm

Initialize:

- $X = [s]$ [vertices processed so far]
- $A[s] = 0$ [computed shortest path distances]
- $B[s] = \text{empty path}$ [computed shortest paths]

Main Loop

- while $X \neq V$:

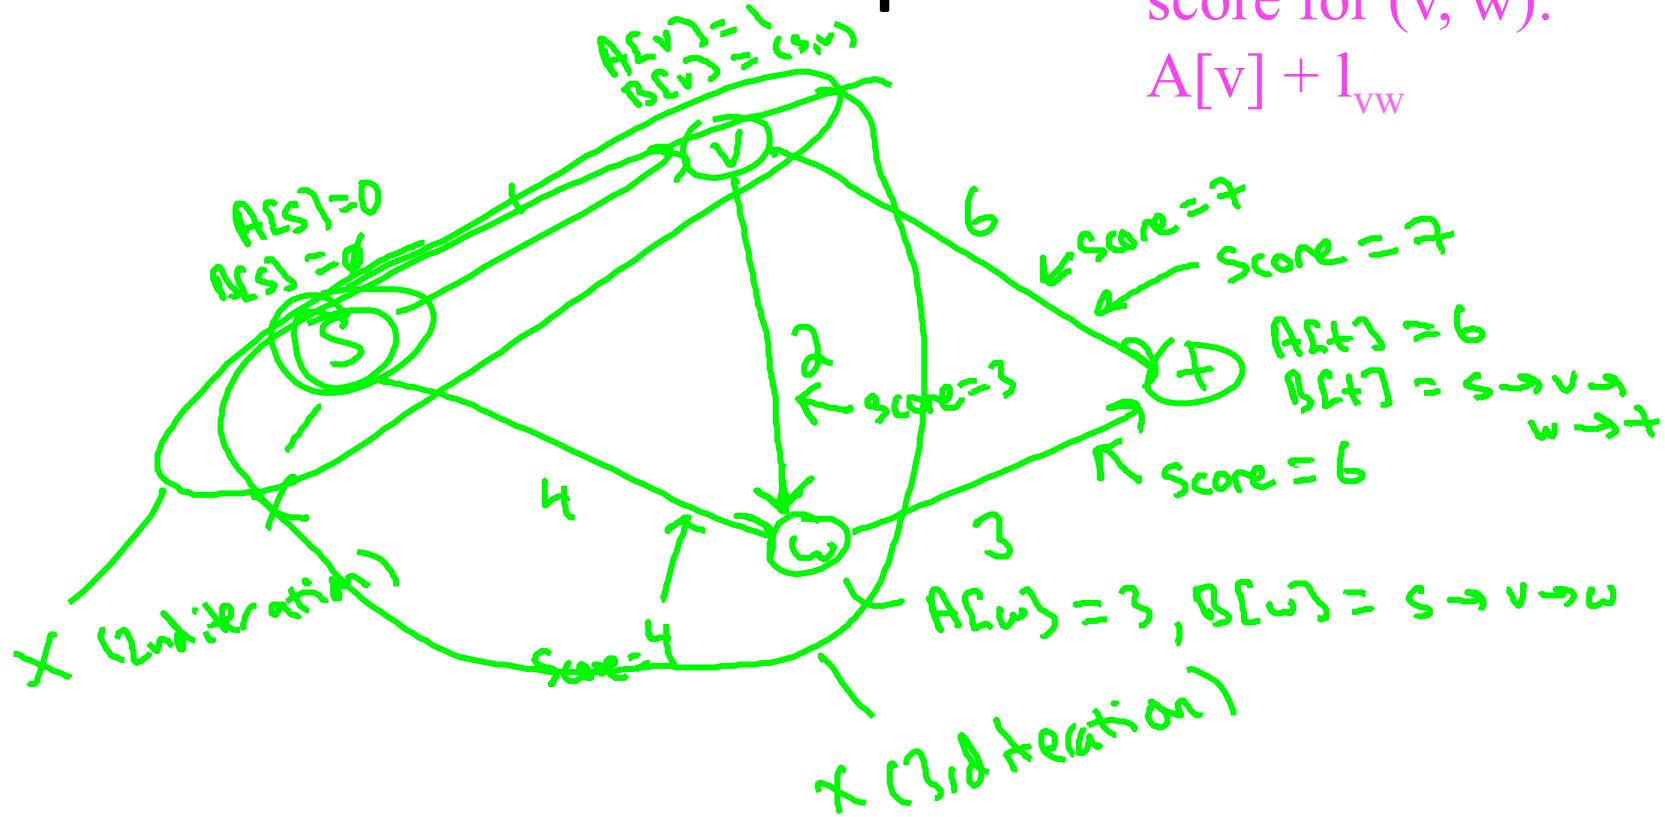


-need to grow
X by one node

Main Loop cont'd:

- among all edges $(v, w) \in E$ with $v \in X, w \notin X$, pick the one that minimizes $A[v] + l_{vw}$
- [call it (v^*, w^*)] Already computed in earlier iteration
- add w^* to X
 - set $A[w^*] := A[v^*] + l_{v^*w^*}$
 - set $B[w^*] := B[v^*] \cup (v^*, w^*)$

Example



Dijkstra's greedy
score for (v, w) :

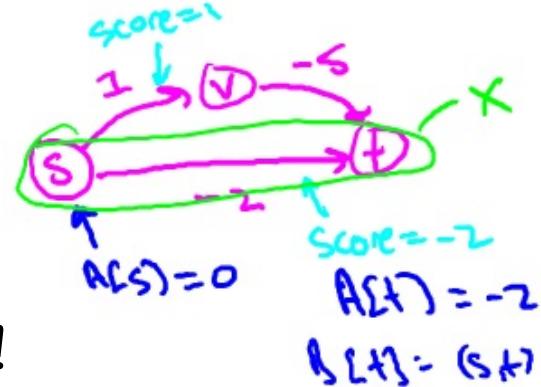
$$A[v] + l_{vw}$$

Non-Example

Question: why not reduce computing shortest paths with negative edge lengths to the same problem with non negative lengths? (by adding large constant to edge lengths)

Problem: doesn't preserve shortest paths !

Also: Dijkstra's algorithm incorrect on this graph !
(computes shortest s-t distance to be -2 rather than -4)





Design and Analysis
of Algorithms I

Graph Primitives

Dijkstra's Algorithm: The Basics

Single-Source Shortest Paths

Input: directed graph $G=(V, E)$. ($m=|E|$, $n=|V|$)

- each edge has non negative length l_e
- source vertex s

Output: for each $v \in V$, compute

$L(v) :=$ length of a shortest $s-v$ path in G

Assumption:

1. [for convenience] $\forall v \in V, \exists s \Rightarrow v$ path
2. [important] $l_e \geq 0 \quad \forall e \in E$

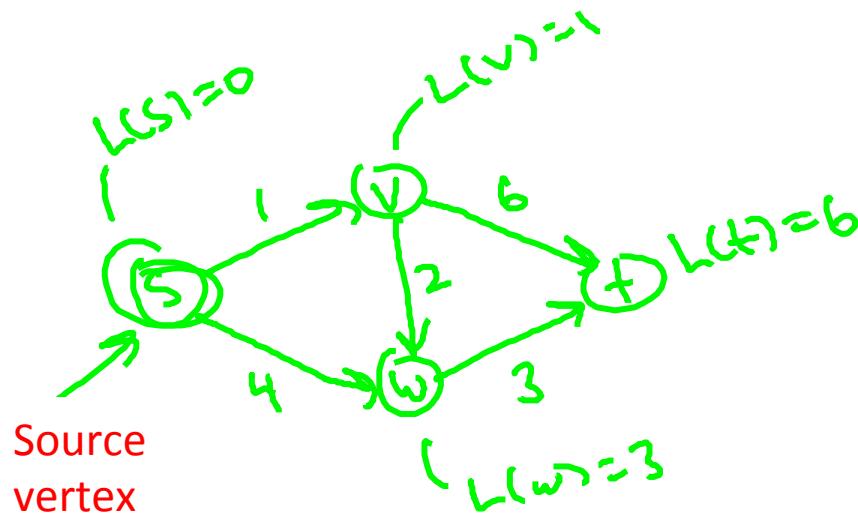
Length of path
= sum of edge lengths



Path length = 6

One of the following is the list of shortest-path distances for the nodes s, v, w, t , respectively. Which is it?

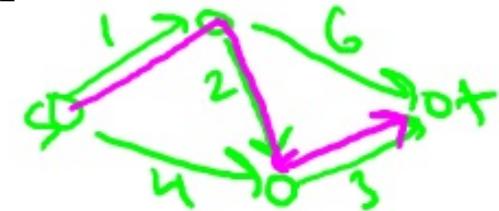
- 0,1,2,3
- 0,1,4,7
- 0,1,4,6
- 0,1,3,6



Why Another Shortest-Path Algorithm?

Question: doesn't BFS already compute shortest paths in linear time?

Answer: yes, IF $l_e = 1$ for every edge e .



Question: why not just replace each edge e by directed path of l_e unit length edges:



Answer: blows up graph too much

Solution: Dijkstra's shortest path algorithm.

This array
only to help
explanation!

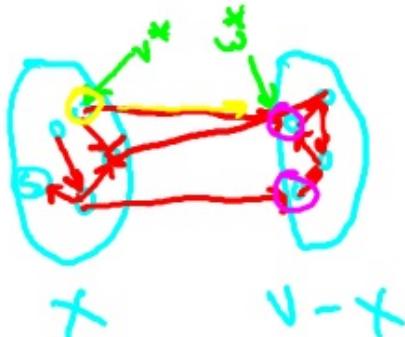
Dijkstra's Algorithm

Initialize:

- $X = [s]$ [vertices processed so far]
- $A[s] = 0$ [computed shortest path distances]
- $B[s] = \text{empty path}$ [computed shortest paths]

Main Loop

- while $X \neq V$:

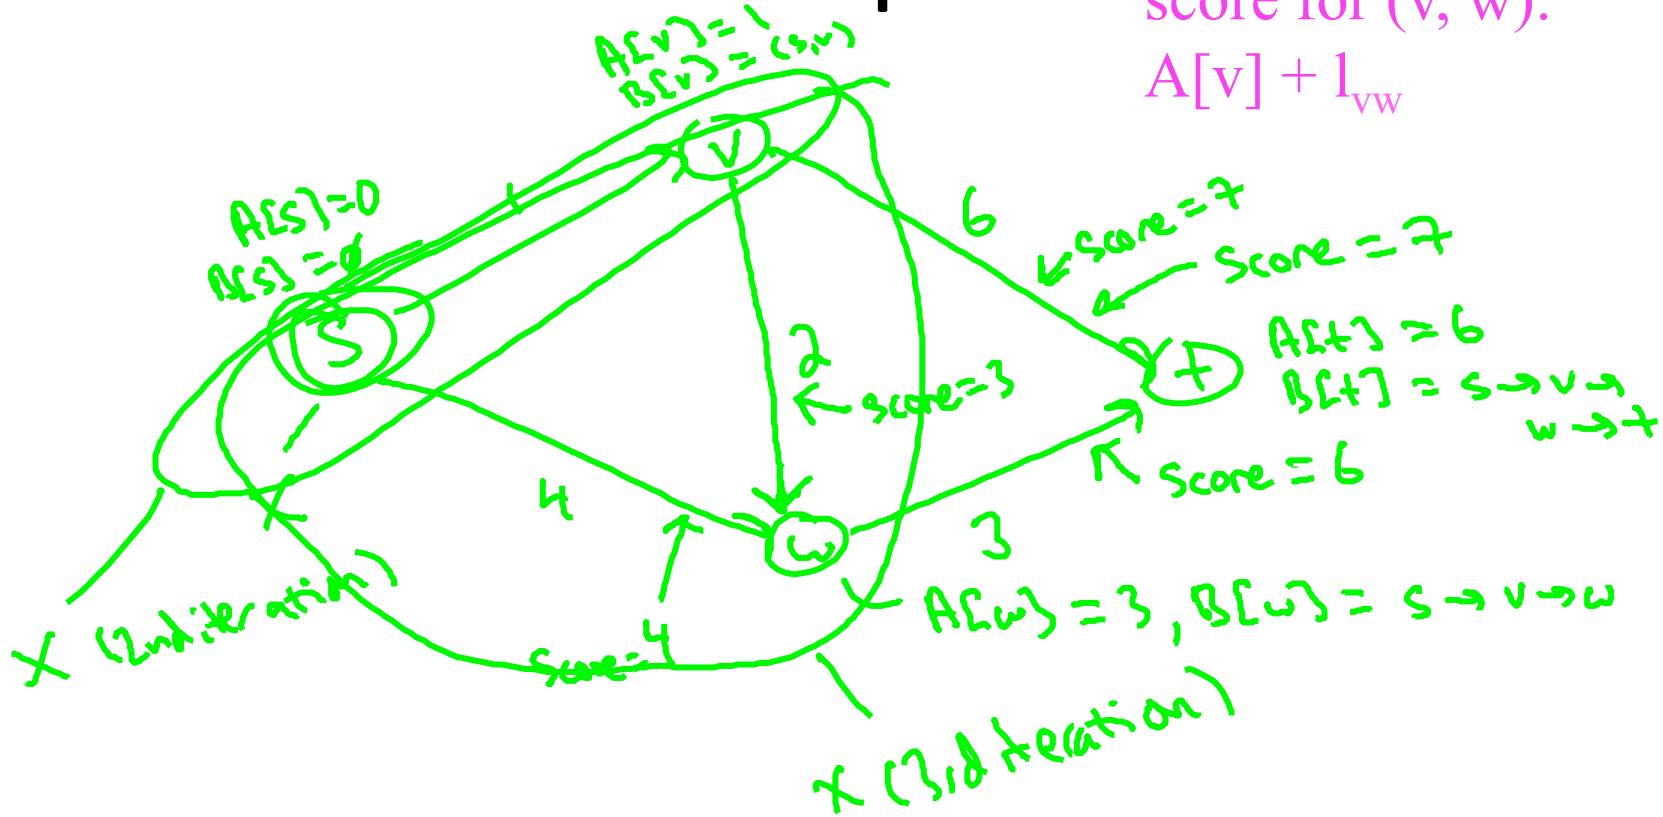


-need to grow
 X by one node

Main Loop cont'd:

- among all edges $(v, w) \in E$ with $v \in X, w \notin X$, pick the one that minimizes $A[v] + l_{vw}$
[call it (v^*, w^*)] Already computed in earlier iteration
- add w^* to X
- set $A[w^*] := A[v^*] + l_{v^*w^*}$
- set $B[w^*] := B[v^*] \cup (v^*, w^*)$

Example



Dijkstra's greedy
score for (v, w) :

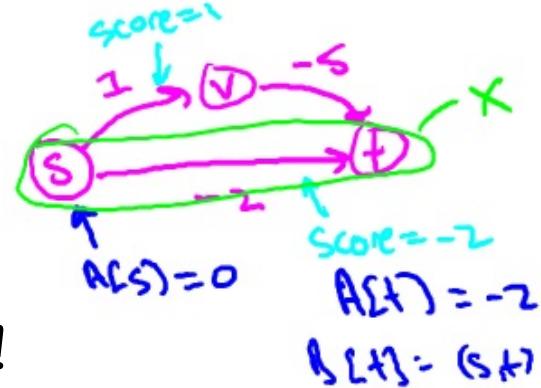
$$A[v] + l_{vw}$$

Non-Example

Question: why not reduce computing shortest paths with negative edge lengths to the same problem with non negative lengths? (by adding large constant to edge lengths)

Problem: doesn't preserve shortest paths !

Also: Dijkstra's algorithm incorrect on this graph !
(computes shortest s-t distance to be -2 rather than -4)





Design and Analysis
of Algorithms I

Graph Primitives

Dijkstra's Algorithm: Why It Works

This array
only to help
explanation!

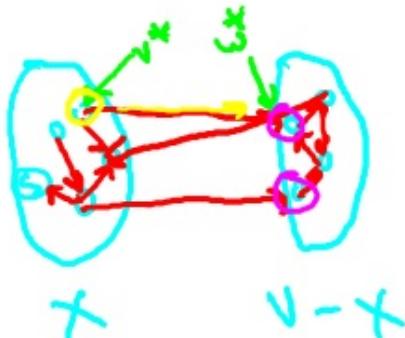
Dijkstra's Algorithm

Initialize:

- $X = [s]$ [vertices processed so far]
- $A[s] = 0$ [computed shortest path distances]
- $B[s] = \text{empty path}$ [computed shortest paths]

Main Loop

- while $X \neq V$:



-need to grow
 X by one node

Main Loop cont'd:

- among all edges $(v, w) \in E$ with $v \in X, w \notin X$, pick the one that minimizes $A[v] + l_{vw}$
[call it (v^*, w^*)] Already computed in earlier iteration
- add w^* to X
- set $A[w^*] := A[v^*] + l_{v^*w^*}$
- set $B[w^*] := B[v^*] \cup (v^*, w^*)$

Correctness Claim

Theorem [Dijkstra] For every directed graph with nonnegative edge lengths, Dijkstra's algorithm correctly computes all shortest-path distances.

$$[i.e., A[v] = L(v) \forall v \in V]$$

what algorithm
computes

True shortest
distance from s to v

Proof: by induction on the number of iterations.

Base Case: $A[s] = L[s] = 0$ (correct)

Proof

Inductive Step:

Inductive Hypothesis: all previous iterations correct (i.e., $A[v] = L(v)$ and $B[v]$ is a true shortest s-v path in G , for all v already in X).

In current iteration:

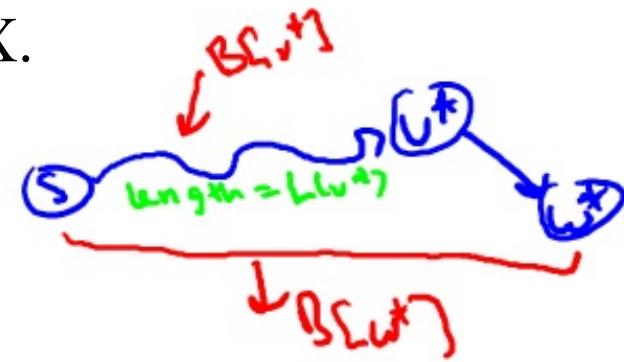
We pick an edge (v^*, w^*) and we add w^* to X .

We set $B[w^*] = B[v^*] \cup (v^*, w^*)$

has length $L(v^*) + l_{v^*w^*}$

has length $L(v^*)$

Also: $A[w^*] = A[v^*] + l_{v^*w^*} = L(v^*) + l_{v^*w^*}$



Proof (con'd)

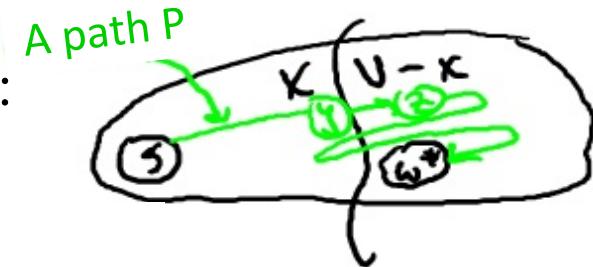
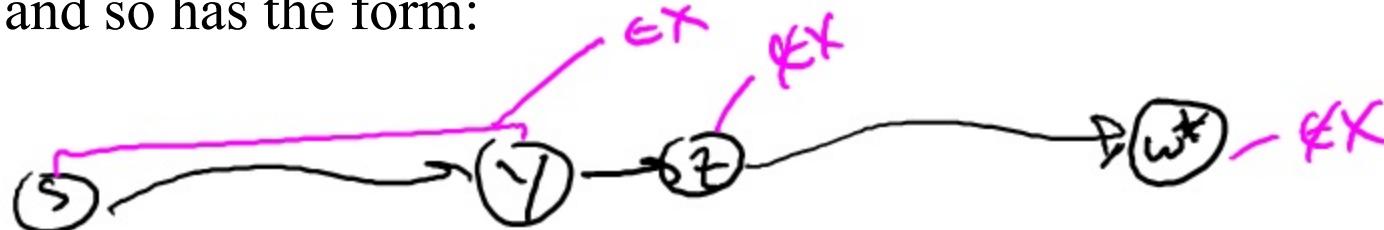
Upshot: in current iteration, we set:

1. $A[w^*] = L(v^*) + l_{v^*w^*}$
2. $B[w^*] = \text{an } s \rightarrow w^* \text{ path with length } (L(v^*) + l_{v^*w^*})$

To finish proof: need to show that *every* $s-w^*$ path has length $\geq L(v^*) + l_{v^*w^*}$ (if so, our path is the shortest!)

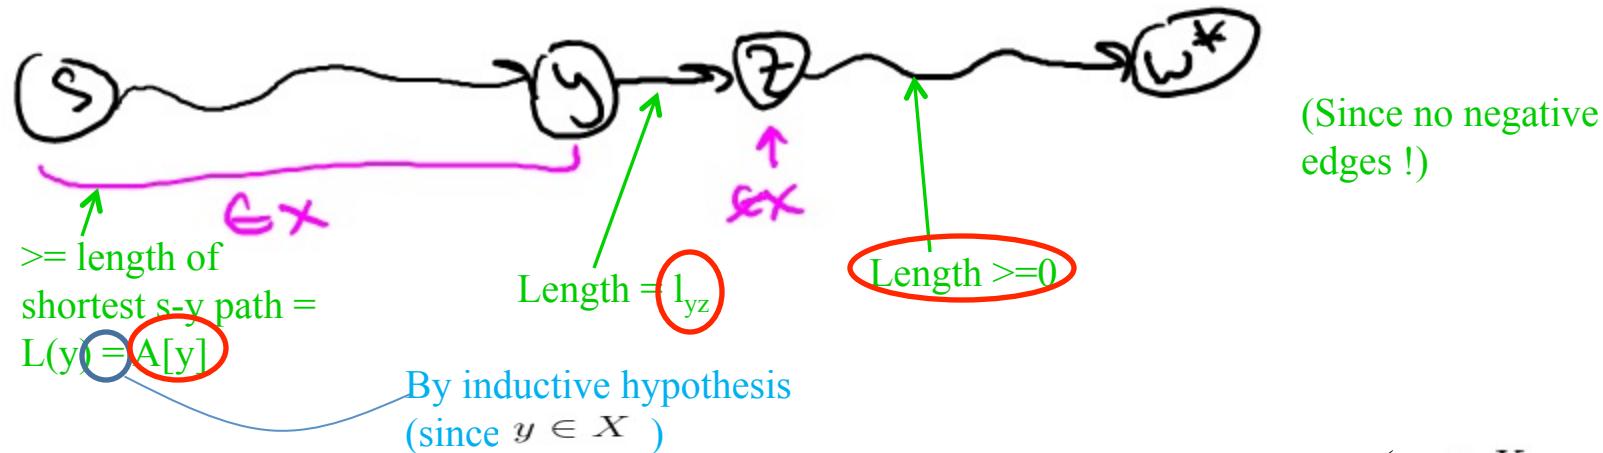
So: Let $P = \text{any } s \rightarrow w^* \text{ path. Must "cross the frontier":}$

and so has the form:



Proof (con'd)

So: every $s \rightarrow w^*$ path P has to have the form



Total length of path P: at least $A[y] + C_{yz}$ ↗ length of our path !
-> by Dijkstra's greedy criterion, $A[v^*] + l_{v^*w^*} \leq A[y] + l_{yz} \leq \text{length of } P$

$(y \in X$
 $z \notin X)$

Q.E.D.



Design and Analysis
of Algorithms I

Graph Primitives

Dijkstra's Algorithm:
Fast Implementation

Single-Source Shortest Paths

Input: directed graph $G=(V, E)$. ($m=|E|$, $n=|V|$)

- each edge has non negative length l_e
- source vertex s

Output: for each $v \in V$, compute

$L(v) :=$ length of a shortest $s-v$ path in G

Assumption:

1. [for convenience] $\forall v \in V, \exists s \Rightarrow v$ path
2. [important] $l_e \geq 0 \quad \forall e \in E$

Length of path
= sum of edge lengths



This array
only to help
explanation!

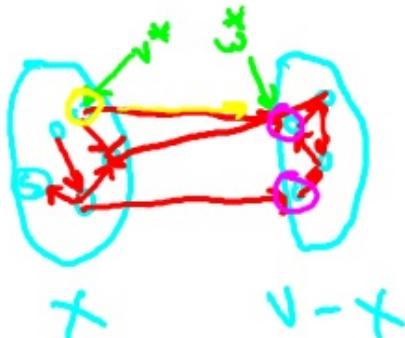
Dijkstra's Algorithm

Initialize:

- $X = [s]$ [vertices processed so far]
- $A[s] = 0$ [computed shortest path distances]
- ~~B[s] = empty path~~ [computed shortest paths]

Main Loop

- while $X \neq V$:



-need to grow
x by one node

Main Loop cont'd:

- among all edges $(v, w) \in E$ with $v \in X, w \notin X$, pick the one that minimizes $A[v] + l_{vw}$
[call it (v^, w^*)]*
- add w^* to X
- set $A[w^*] := A[v^*] + l_{v^*w^*}$
- set ~~$B[w^*] := B[v^*] \cup (v^*, w^*)$~~

Already
computed in
earlier iteration

Which of the following running times seems to best describe a “naïve” implementation of Dijkstra’s algorithm?

- $\theta(m+n)$
- $\theta(m \log n)$
- $\theta(n^{12})$
- $\theta(mn)$

- (n-1) iterations of while loop
- $\theta(m)$ work per iteration
 - [$\theta(1)$ work per edge]

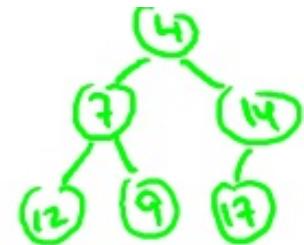
CAN WE DO BETTER?

Heap Operations

Recall: raison d'être of heap = perform Insert, Extract-Min in $O(\log n)$ time.
[rest of video assumes familiarity with heaps]

Height $\sim \log_2 n$

- conceptually, a perfectly balanced binary tree
- Heap property: at every node, key \leq children's keys
- extract-min by swapping up last leaf, bubbling down
- insert via bubbling up



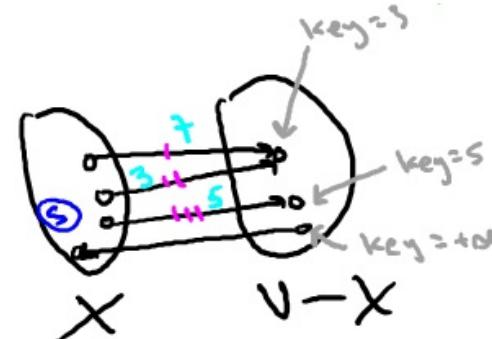
Also: will need ability to delete from middle of heap. (bubble up or down as needed)

Two Invariants

Invariant #1: elements in heap = vertices of $V-X$.

Invariant #2: for $v \notin X$
 $\text{Key}[v] = \text{smallest Dijkstra greedy score of an edge } (u, v) \text{ in } E \text{ with } v \in X$

(of $+\infty$ if no such edges exist)



Point: by invariants, Extract-Min yields correct vertex w^* to add to X next.

(and we set $A[w^*]$ to $\text{key}[w^*]$)

Dijkstra's
greedy score
of (v, w) :
 $A[v] + l_{vw}$

Maintaining the Invariants

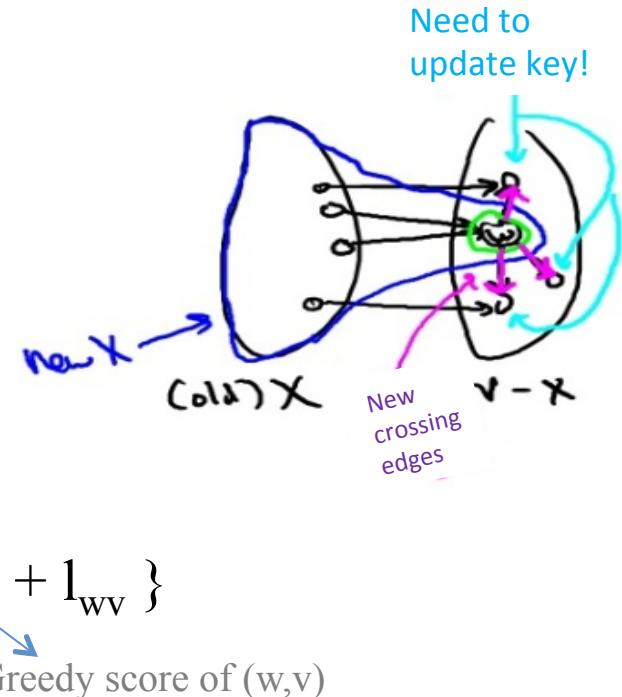
To maintain Invariant #2: [i.e., that $\forall v \notin X$

$\text{Key}[v] = \text{smallest Dijkstra greedy score of edge } (u,v) \text{ with } u \text{ in } X$]

When w extracted from heap (i.e., added to X)

- for each edge (w,v) in E:
 - if $v \in V-X$ (i.e., in heap)
 - delete v from heap
 - recompute $\text{key}[v] = \min \{\text{key}[v], A[w] + l_{wv} \}$
 - re-Insert v into heap

Key update



Running Time Analysis

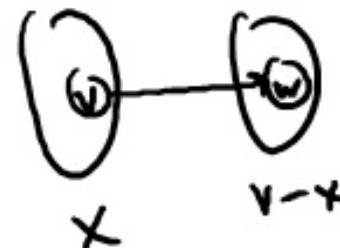
You check: dominated by heap operations. ($O(\log(n))$ each)

- $(n-1)$ Extract mins
- each edge (v,w) triggers at most one Delete/Insert combo
(if v added to X first)

So: # of heap operations in $O(n+m) = O(m)$

So: running time = $O(m \log(n))$ (like sorting)

Since graph is
weakly connected





Data Structures

Introduction

Design and Analysis
of Algorithms I

Data Structures

Point : organize data so that it can be accessed quickly and usefully.

Examples : lists, stacks, queues, heaps, search trees, hashtables, bloom filters, union-find, etc.

Why so Many ? : different data structures support different sets of operations => suitable for different types of tasks.

Rule of Thumb : choose the “minimal” data structure that supports all the operations that you need.

Taking It To The Next Level

Level 0

- “what’s a data structure ?”

Level 1

- cocktail party-level literacy

Level 2

- “this problem calls out for a heap”

Level 3

- “I only use data structures that I create myself”





Design and Analysis
of Algorithms I

Data Structures

Heaps and Their Applications

Heap: Supported Operations

- A container for objects that have keys
- Employer records, network edges, events, etc.

Insert: add a new object to a heap.

Running time : $O(\log(n))$

Equally well,
EXTRACT MAX

Extract-Min: remove an object in heap with a minimum key value. [ties broken arbitrarily]

Running time : $O(\log n)$ [$n = \#$ of objects in heap]

Also : **HEAPIFY** ($\binom{n \text{ batched Inserts}}{\text{in } O(n) \text{ time}}$), **DELETE**($O(\log(n))$ time)

Application: Sorting

Canonical use of heap : fast way to do repeated minimum computations.

Example : SelectionSort $\sim \theta(n)$ linear scans, $\theta(n^2)$ runtime on array of length n

Heap Sort : 1.) insert all n array elements into a heap
2.) Extract-Min to pluck out elements in sorted order

Running Time = $2n$ heap operations = $O(n \log(n))$ time.

=> optimal for a “comparison-based” sorting algorithm!

Application: Event Manager

“Priority Queue” – synonym for a heap.

Example : simulation (e.g., for a video game)

- Objects = event records [Action/update to occur at given time in the future]
- Key = time event scheduled to occur
- Extract-Min => yields the next scheduled event

Application: Median Maintenance

I give you : a sequence x_1, \dots, x_n of numbers, one-by-one.

You tell me : at each time step i , the median of $\{x_1, \dots, x_i\}$.

Constraint : use $O(\log(i))$ time at each step i .

Solution : maintain heaps H_{Low} : supports Extract Max

H_{High} : supports Extract Min

Key Idea : maintain invariant that $\sim i/2$ smallest (largest) elements in
 $H_{\text{Low}} (H_{\text{High}})$

You Check : 1.) can maintain invariant with $O(\log(i))$ work
2.) given invariant, can compute median in $O(\log(i))$ work

Application: Speeding Up Dijkstra

Dijkstra's Shortest-Path Algorithm

- Naïve implementation => runtime = $O(nm)$
- with heaps => runtime = $O(m \log(n))$

$$\theta(nm)$$

vertices # edges
 ↓
 ↓
loop iterations Work per iteration
 [linear scan through
 edges for minimum
 computation]



Design and Analysis
of Algorithms I

Data Structures

Heaps: Some
Implementation Details

Heap: Supported Operations

- A container for objects that have keys
- Employer records, network edges, events, etc.

Insert: add a new object to a heap.

Running time : $O(\log(n))$

Equally well,
EXTRACT MAX

Extract-Min: remove an object in heap with a minimum key value. [ties broken arbitrarily]

Running time : $O(\log n)$ [$n = \#$ of objects in heap]

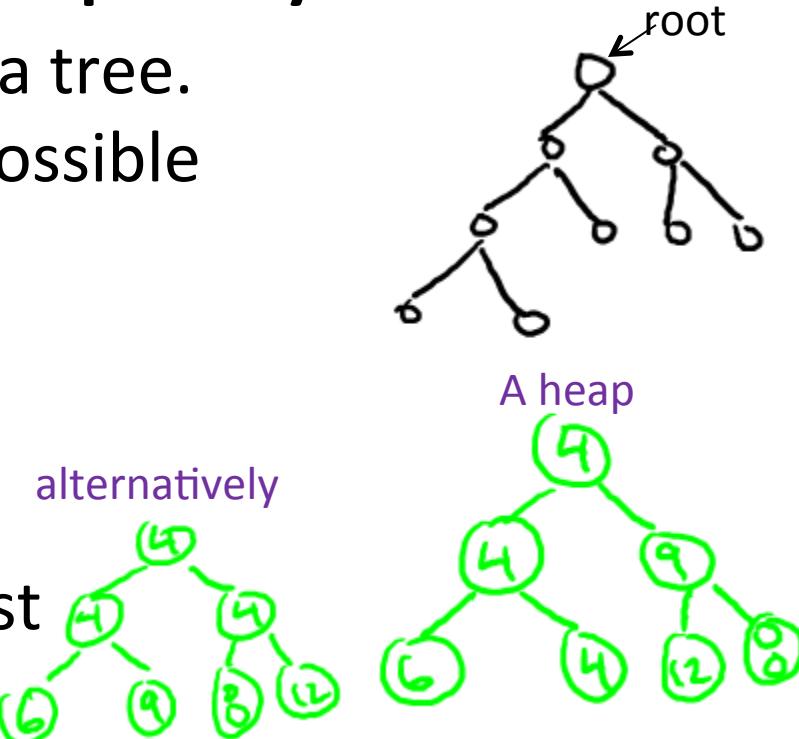
Also : **HEAPIFY** ($\binom{n \text{ batched Inserts}}{\text{in } O(n) \text{ time}}$), **DELETE**($O(\log(n))$ time)

The Heap Property

Conceptually : think of a heap as a tree.
-rooted, binary, as complete as possible

Heap Property: at every node x ,
 $\text{Key}[x] \leq \text{all keys of } x\text{'s children}$

Consequence : object at root must
have minimum key value

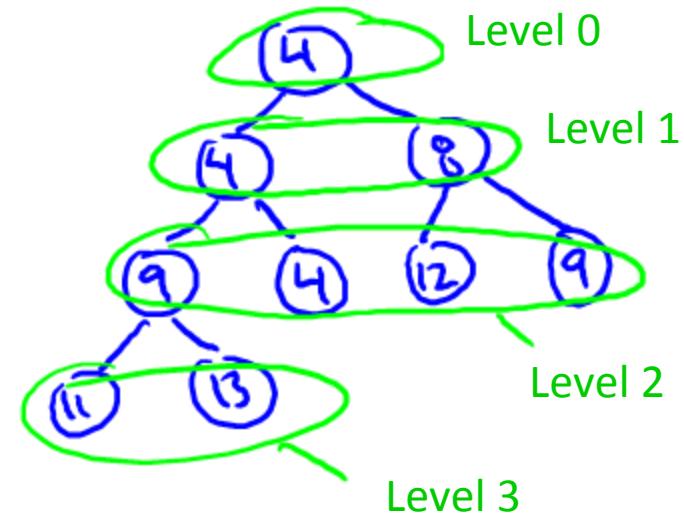


Array Implementation



Note : $\text{parent}(i) = i/2$ if i even
= $[i/2]$ if i odd

i.e., round down



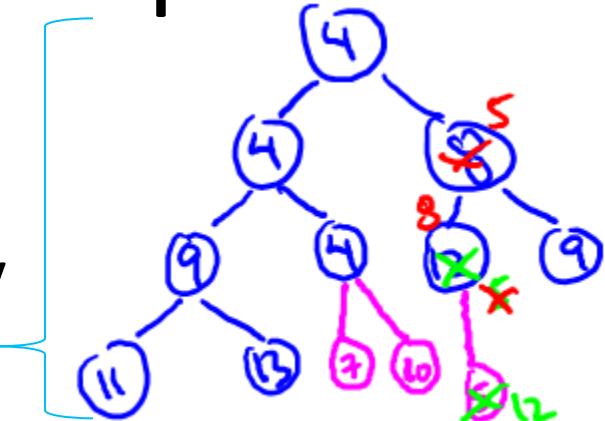
and children of i are $2i, 2i+1$

Insert and Bubble-Up

Implementation of Insert (given key k)

Step 1: stick k at end of last level.

Step 2 : Bubble-Up k until heap property is restored (i.e., key of k's parent is $\leq k$)



$\sim \log_2 n$ levels ($n = \#$ of items in heap)

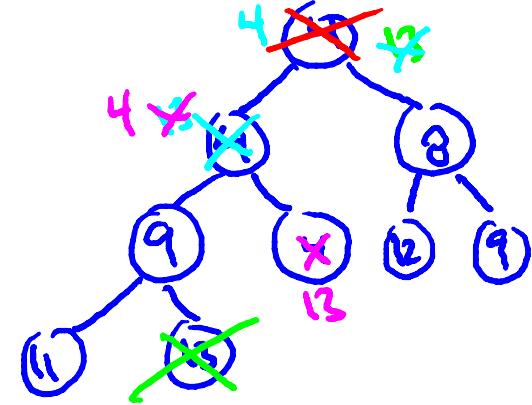
Check : 1.) bubbling up process must stop, with heap property restored
2.) runtime = $O(\log(n))$

Extract-Min and Bubble-Down

Implementation of Extract-Min

1. Delete root
2. Move last leaf to be new root.
3. Iteratively Bubble-Down until heap property has been restored

[always swap with smaller child!]



Check : 1.) only Bubble-Down once per level, halt with a heap
2.) run time = $O(\log(n))$

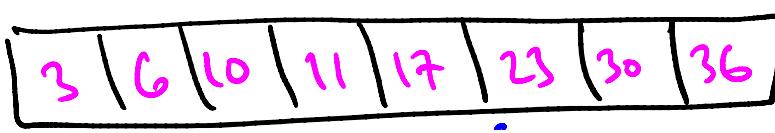


Design and Analysis
of Algorithms I

Data Structures

Balanced Search
Trees: Supported
Operations

Sorted Arrays: Supported Operations



OPERATIONS

SEARCH

SELECT (given order statistic I)

MIN/MAX

PRED/SUCC (given pointer to a key)

RANK (i.e., # of keys less than or equal to
a given value)

OUTPUT IN SORTED ORDER

BUT WHAT ABOUT
INSERTIONS + DELETIONS ?
(would take $\Theta(n)$ time)

RUNNING TIME

$\Theta(\log(n))$

$O(1)$

$O(1)$

$O(1)$

$\Theta(\log(n))$

$O(n)$

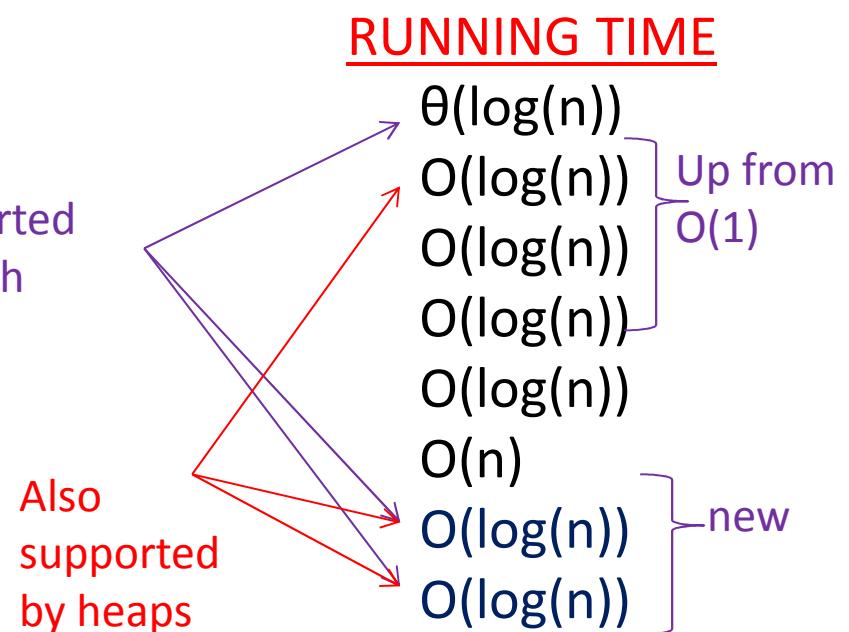
Balanced Search Trees: Supported Operations

Raison d'etre : like sorted array + fast (logarithmic) inserts + deletes !

OPERATIONS

SEARCH
SELECT
MIN/MAX
PRED/SUCC
RANK
OUTPUT IN SORTED ORDER
INSERT
DELETE

Also supported by hash tables



Tim Roughgarden



Design and Analysis
of Algorithms I

Data Structures

Binary Search
Tree Basics

Balanced Search Trees: Supported Operations

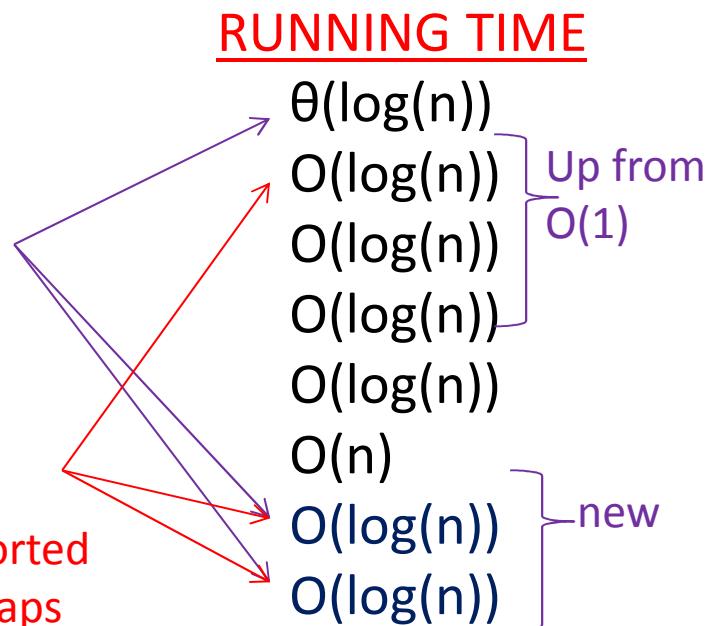
Raison d'etre : like sorted array + fast (logarithmic) inserts + deletes !

OPERATIONS

SEARCH
SELECT
MIN/MAX
PRED/SUCC
RANK
OUTPUT IN SORTED ORDER
INSERT
DELETE

Also supported by hash tables

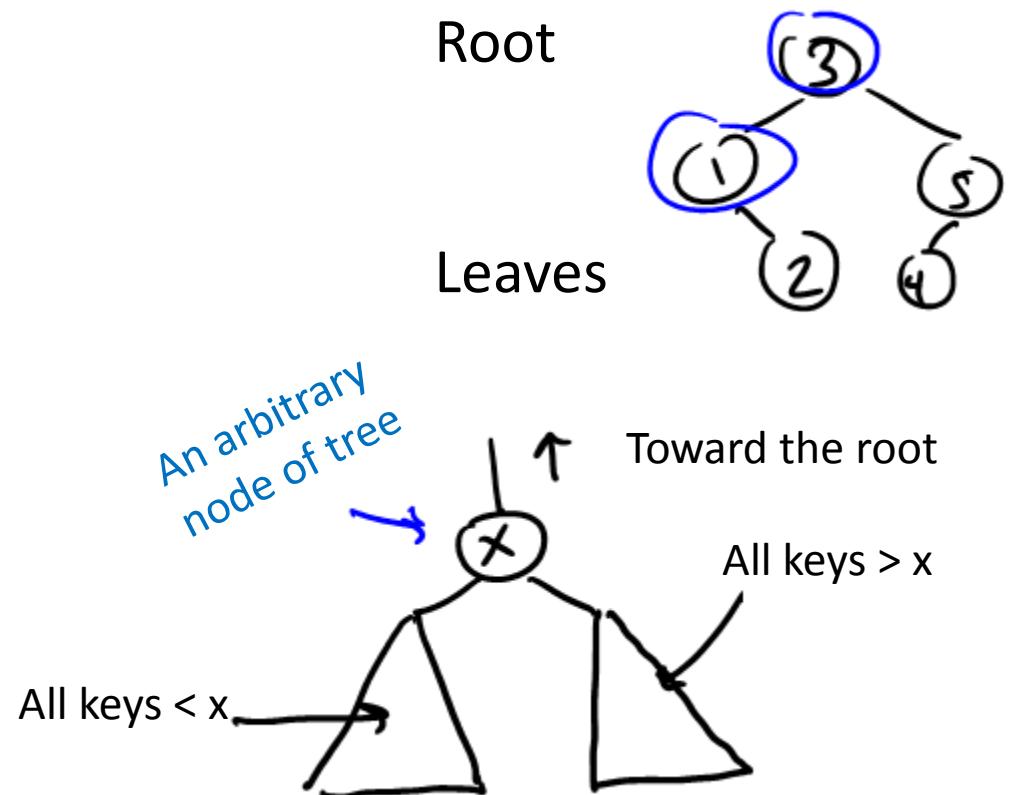
Also supported by heaps



Binary Search Tree Structure

- exactly one node per key
- most basic version :
 - each node has
 - left child pointer
 - right child pointer
 - parent pointer

SEARCH TREE PROPERTY :
(should hold at every node of the search tree)



Tim Roughgarden

The Height of a BST

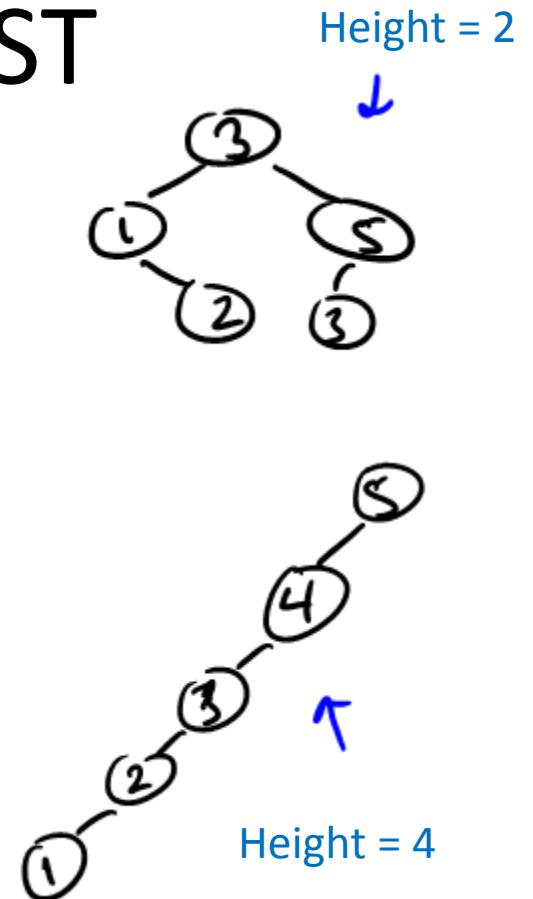
Note : many possible trees for a set of keys.

Note : height could be anywhere from $\sim \log_2 n$ to $\sim n$

(aka depth) longest root-leaf path

Worst case, a chain

Best case, perfectly balanced



Tim Roughgarden

Searching and Inserting

To Search for key k in tree T

- start at the root
- traverse left / right child pointers as needed

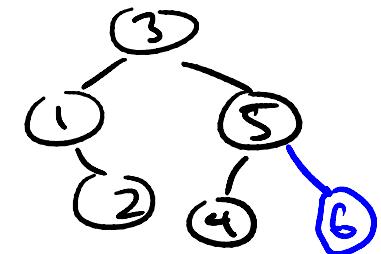
If $k <$ key at current node If $k >$ key at current node

- return node with key k or NULL, as appropriate

To Insert a new key k into a tree T

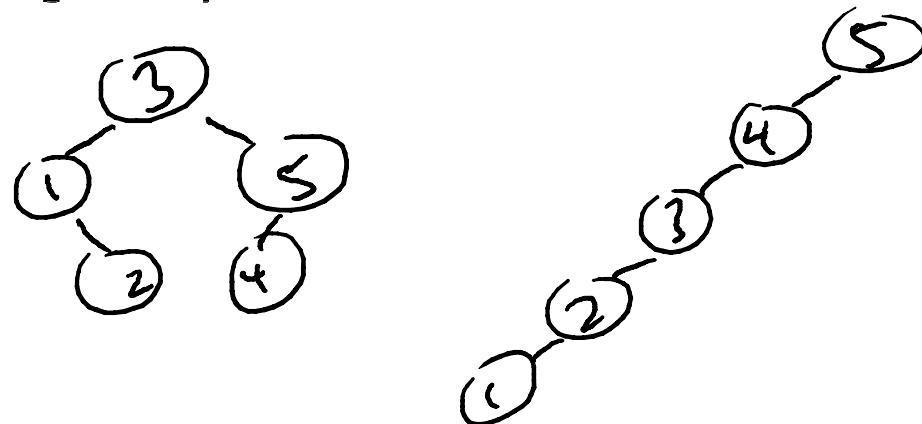
- search for k (unsuccessfully)
- rewire final NULL ptr to point to new node with key k

Exercise :
preserves
search tree
property!



The worst-case running time of Search (or Insert) operation in a binary search tree containing n keys is...?

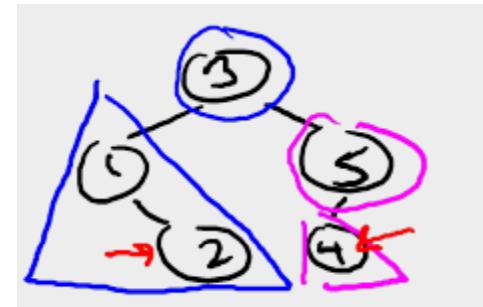
- $\theta(1)$
- $\theta(\log_2 n)$
- $\theta(\text{height})$
- $\theta(n)$



Min, Max, Pred, And Succ

To compute the minimum (maximum) key of a tree

- Start at root
- Follow left child pointers (right ptrs, for maximum) until you can't anymore (return last key found)



To compute the predecessor of key k

- Easy case : If k's left subtree nonempty, return max key in left subtree
- Otherwise : follow parent pointers until you get to a key less than k.

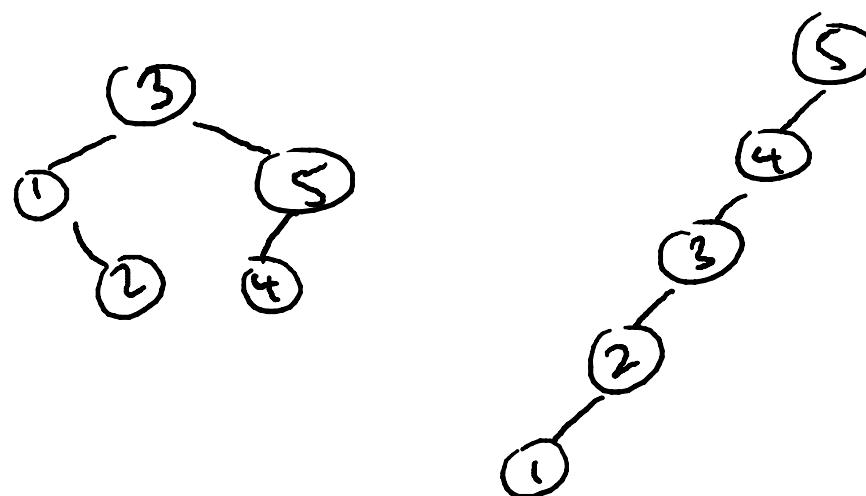
Happens first time you “turn left”

Exercise :
prove this
works

Tim Roughgarden

The worst-case running time of the Max operation in a binary search tree containing n keys is...?

- $\theta(1)$
- $\theta(\log_2 n)$
- $\theta(\text{height})$
- $\theta(n)$



In-Order Traversal

TO PRINT OUT KEYS IN INCREASING ORDER

-Let r = root of search tree, with subtrees TL and TR

- recurse on TL

[by recursion (induction) prints out keys of TL
in increasing order]

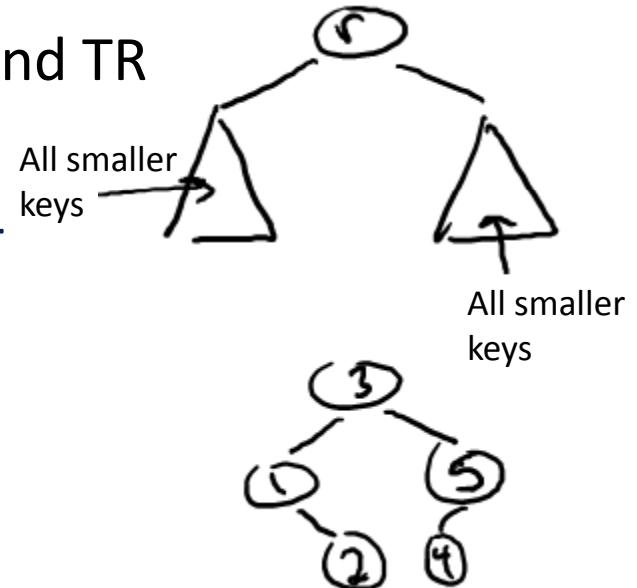
-Print out r 's key

RUNNING TIME

$O(1)$ time, n recursive
calls $\Rightarrow O(n)$ total

-Recurse on TR

[prints out keys of TR in increasing order]



Tim Roughgarden

Deletion

TO DELETE A KEY K FROM A SEARCH TREE

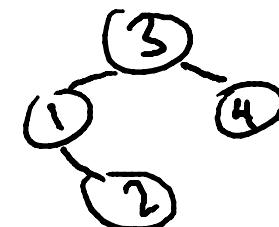
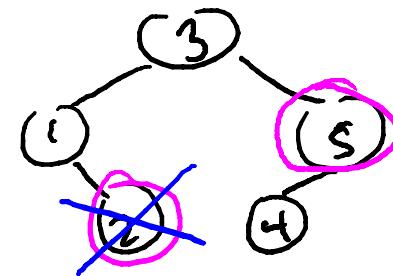
- SEARCH for k

EASY CASE (k's node has no children)

- Just delete k's node from tree, done

MEDIUM CASE (k's node has one child)

(unique child assumes position
previously held by k's node)



Deletion (con'd)

DIFFICULT CASE (k's node has 2 children)

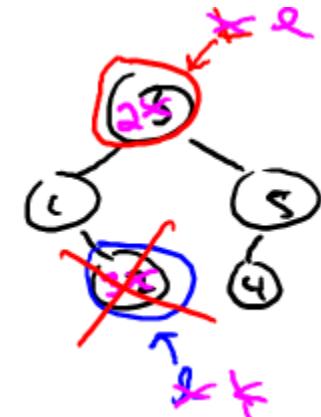
-Compute k's predecessor l

[i.e., traverse k's (non-NULL) left child ptr, then right child ptrs until no longer possible]

- SWAP k and l !

NOTE : in it's new position, k has no right child !

=> easy to delete or splice out k's new node



RUNNING

TIME :

$\Theta(\text{height})$

Exercise : at end, have a valid search tree !

Select and Rank

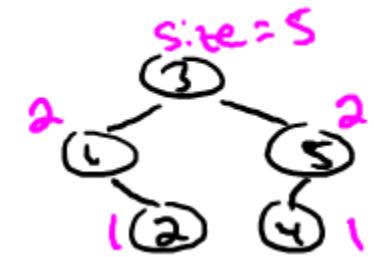
Idea : store a little bit of extra info at each tree node about the tree itself (i.e., not about the data)

Example Augmentation : $\text{size}(x) = \# \text{ of tree nodes in subtree rooted at } x$.

Note : if x has children y and z ,
then $\text{size}(y) + \text{size}(z) + 1$

Population in left subtree Right subtree x itself

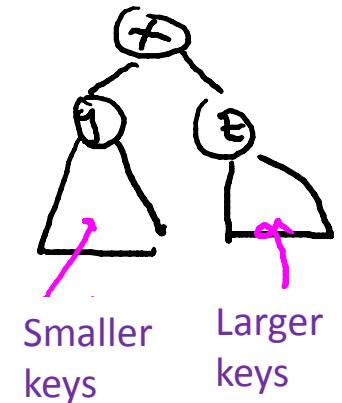
Also : easy to keep sizes up-to-date during an Insertion or Deletion (you check!)



Select and Rank (con'd)

HOW TO SELECT i^{th} ORDER STATISTIC FROM AUGMENTED SEARCH TREE (with subtree sizes)

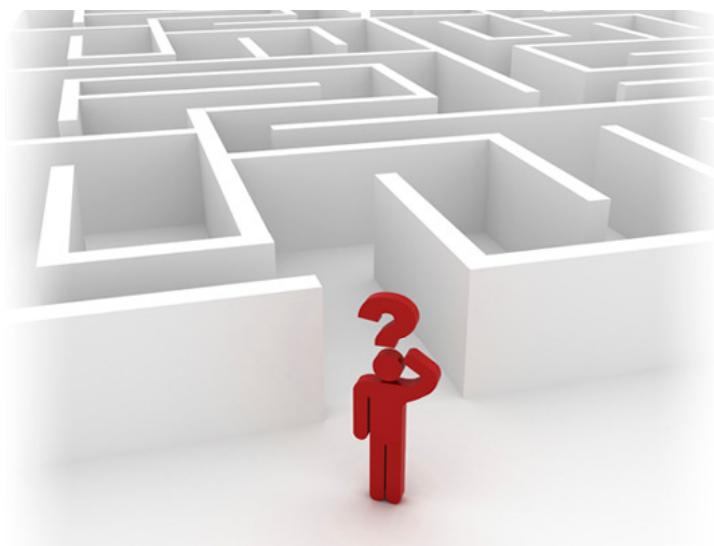
- start at root x , with children y and z
- let $a = \text{size}(y)$ [$a = 0$ if x has no left child]
- if $a = i-1$, return x 's key
- if $a \geq i$, recursively compute i^{th} order statistic of search tree rooted at y
- if $a < i-1$ recursively compute $(i-a-1)^{\text{th}}$ order statistic of search tree rooted at z



RUNNING TIME = $\Theta(\text{height})$.

[EXERCISE : how to implement RANK ?

Tim Roughgarden



Design and Analysis
of Algorithms I

Data Structures

Binary Search
Tree Basics

Balanced Search Trees: Supported Operations

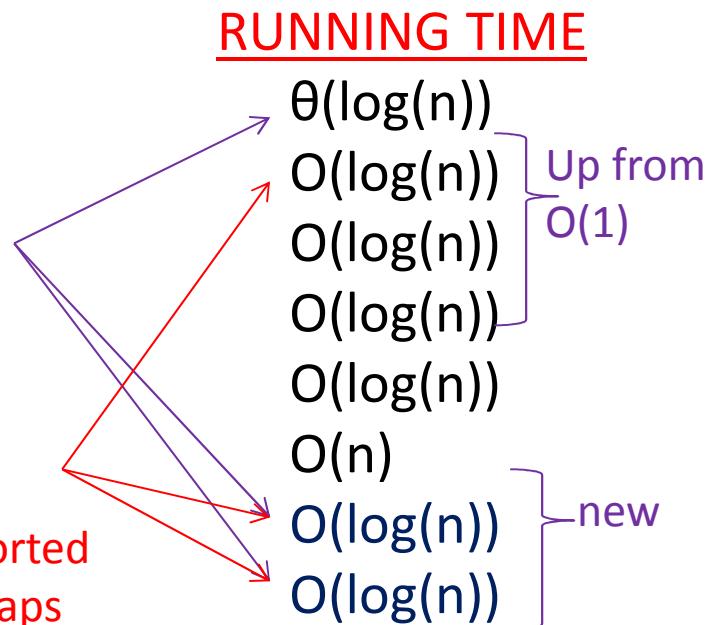
Raison d'etre : like sorted array + fast (logarithmic) inserts + deletes !

OPERATIONS

SEARCH
SELECT
MIN/MAX
PRED/SUCC
RANK
OUTPUT IN SORTED ORDER
INSERT
DELETE

Also supported by hash tables

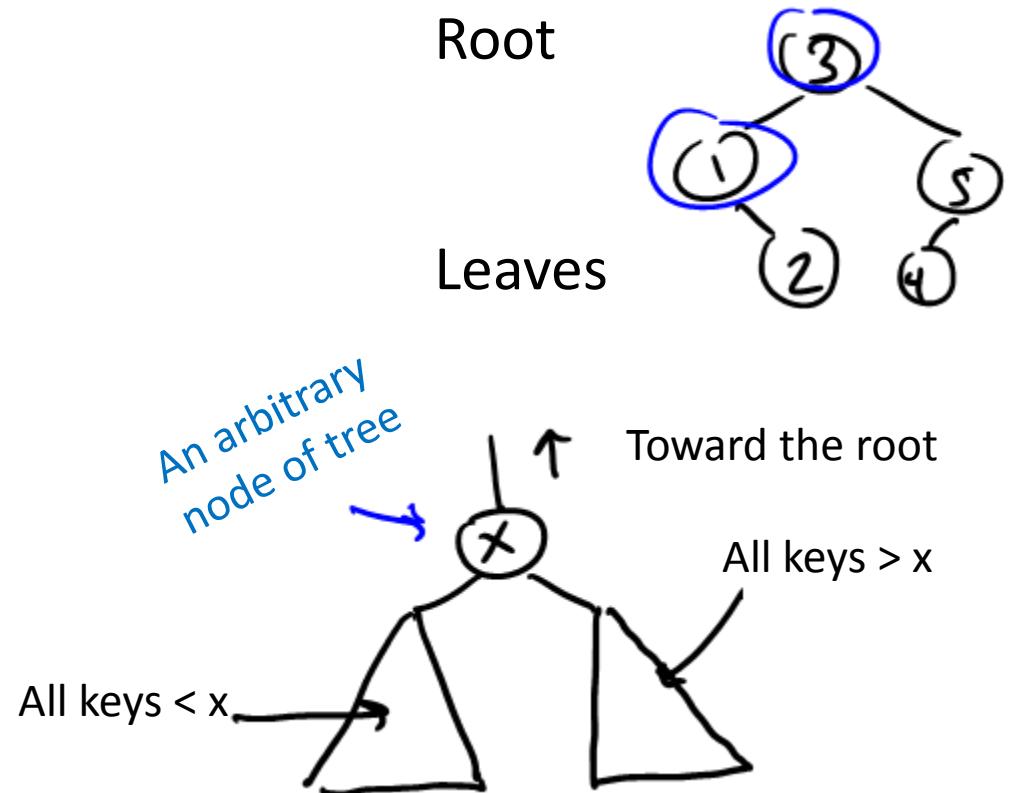
Also supported by heaps



Binary Search Tree Structure

- exactly one node per key
- most basic version :
 - each node has
 - left child pointer
 - right child pointer
 - parent pointer

SEARCH TREE PROPERTY :
(should hold at every node of the search tree)



Tim Roughgarden

The Height of a BST

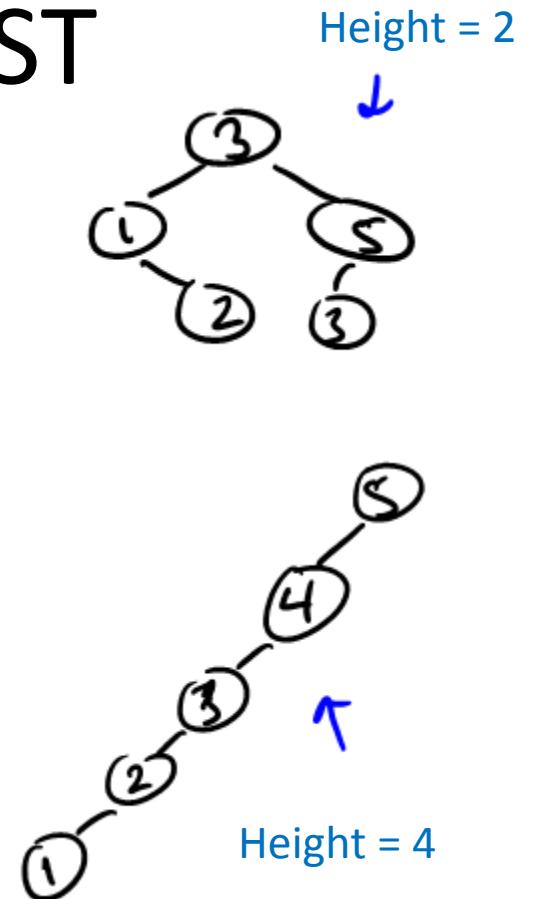
Note : many possible trees for a set of keys.

Note : height could be anywhere from $\sim \log_2 n$ to $\sim n$

(aka depth) longest root-leaf path

Worst case, a chain

Best case, perfectly balanced



Tim Roughgarden

Searching and Inserting

To Search for key k in tree T

- start at the root
- traverse left / right child pointers as needed

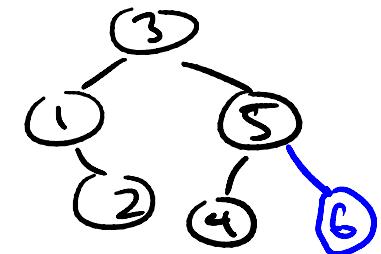
If $k <$ key at current node If $k >$ key at current node

- return node with key k or NULL, as appropriate

To Insert a new key k into a tree T

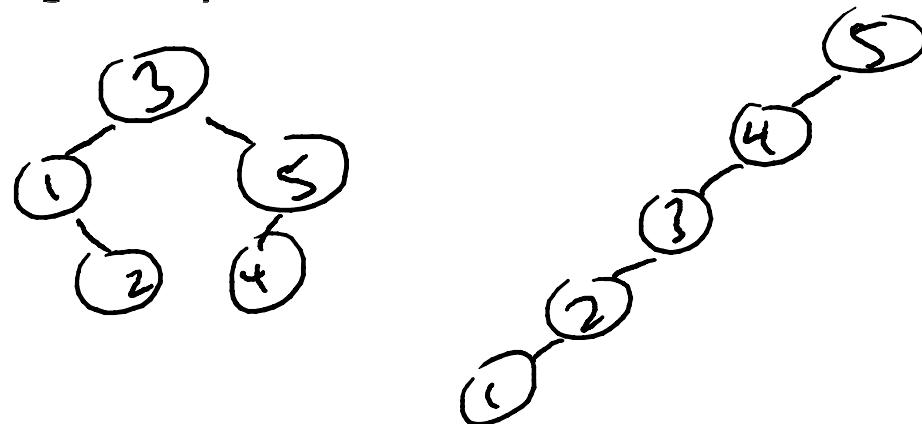
- search for k (unsuccessfully)
- rewire final NULL ptr to point to new node with key k

Exercise :
preserves
search tree
property!



The worst-case running time of Search (or Insert) operation in a binary search tree containing n keys is...?

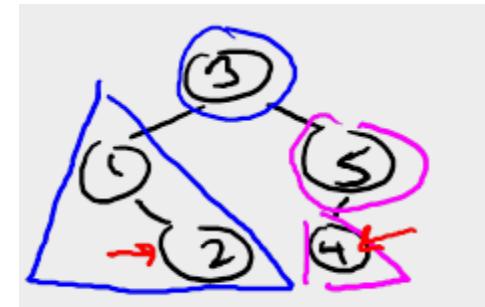
- $\theta(1)$
- $\theta(\log_2 n)$
- $\theta(\text{height})$
- $\theta(n)$



Min, Max, Pred, And Succ

To compute the minimum (maximum) key of a tree

- Start at root
- Follow left child pointers (right ptrs, for maximum) until you can't anymore (return last key found)



To compute the predecessor of key k

- Easy case : If k's left subtree nonempty, return max key in left subtree
- Otherwise : follow parent pointers until you get to a key less than k.

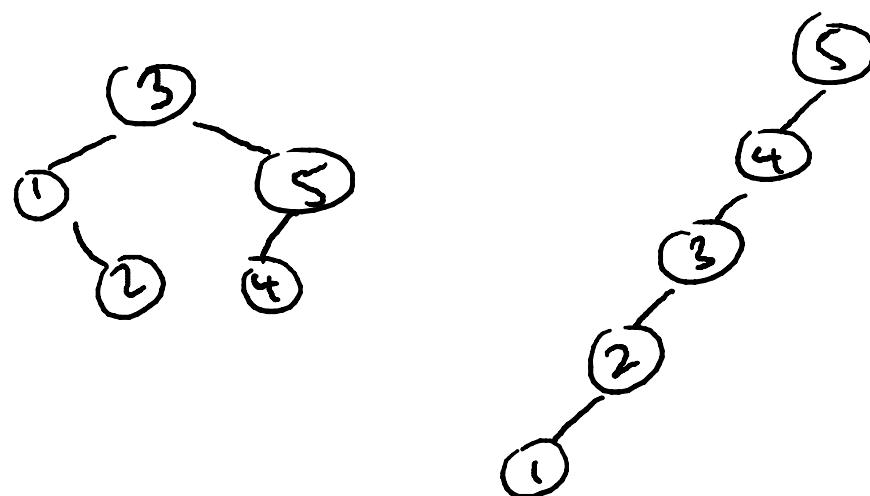
Happens first time you “turn left”

Exercise :
prove this
works

Tim Roughgarden

The worst-case running time of the Max operation in a binary search tree containing n keys is...?

- $\theta(1)$
- $\theta(\log_2 n)$
- $\theta(\text{height})$
- $\theta(n)$



In-Order Traversal

TO PRINT OUT KEYS IN INCREASING ORDER

-Let r = root of search tree, with subtrees TL and TR

- recurse on TL

[by recursion (induction) prints out keys of TL
in increasing order]

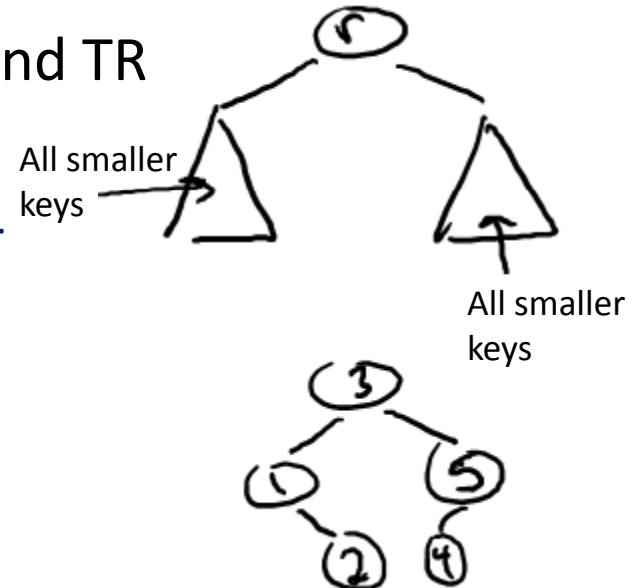
-Print out r 's key

RUNNING TIME

$O(1)$ time, n recursive
calls $\Rightarrow O(n)$ total

-Recurse on TR

[prints out keys of TR in increasing order]



Tim Roughgarden

Deletion

TO DELETE A KEY K FROM A SEARCH TREE

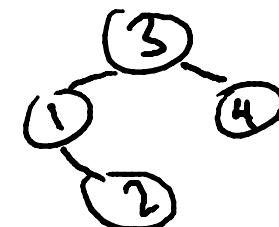
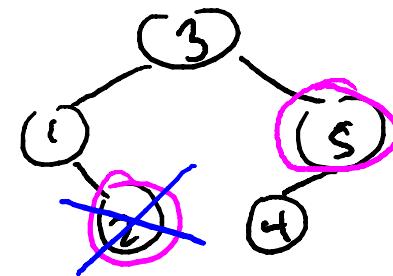
- SEARCH for k

EASY CASE (k's node has no children)

-Just delete k's node from tree, done

MEDIUM CASE (k's node has one child)

(unique child assumes position
previously held by k's node)



Deletion (con'd)

DIFFICULT CASE (k's node has 2 children)

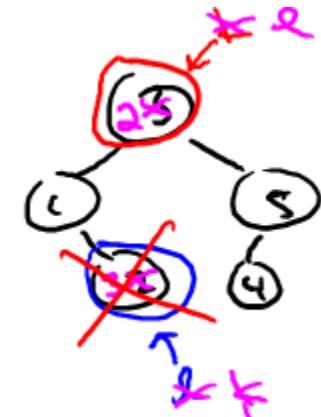
-Compute k's predecessor l

[i.e., traverse k's (non-NULL) left child ptr, then right child ptrs until no longer possible]

- SWAP k and l !

NOTE : in it's new position, k has no right child !

=> easy to delete or splice out k's new node



RUNNING

TIME :

$\theta(\text{height})$

Exercise : at end, have a valid search tree !

Select and Rank

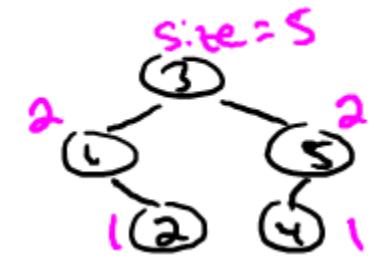
Idea : store a little bit of extra info at each tree node about the tree itself (i.e., not about the data)

Example Augmentation : $\text{size}(x) = \# \text{ of tree nodes in subtree rooted at } x$.

Note : if x has children y and z ,
then $\text{size}(y) + \text{size}(z) + 1$

Population in left subtree Right subtree x itself

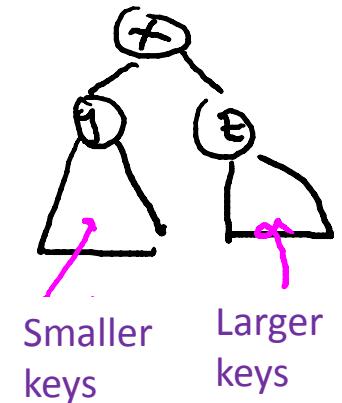
Also : easy to keep sizes up-to-date during an Insertion or Deletion (you check!)



Select and Rank (con'd)

HOW TO SELECT i^{th} ORDER STATISTIC FROM AUGMENTED SEARCH TREE (with subtree sizes)

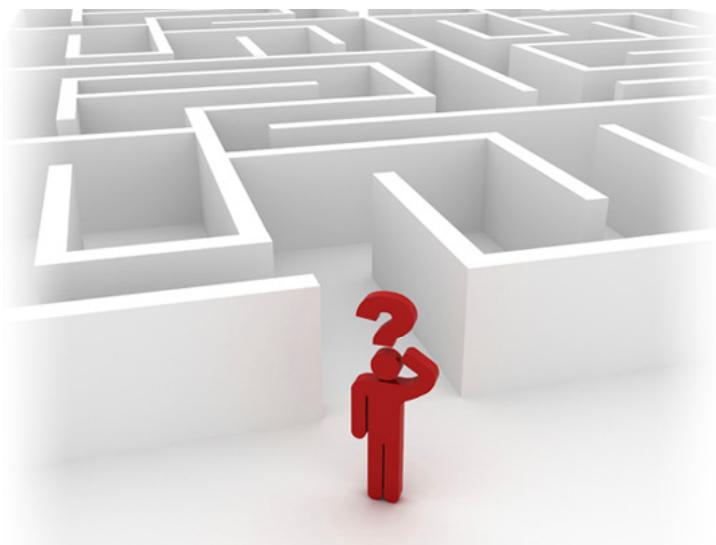
- start at root x , with children y and z
- let $a = \text{size}(y)$ [$a = 0$ if x has no left child]
- if $a = i-1$, return x 's key
- if $a \geq i$, recursively compute i^{th} order statistic of search tree rooted at y
- if $a < i-1$ recursively compute $(i-a-1)^{\text{th}}$ order statistic of search tree rooted at z



RUNNING TIME = $\Theta(\text{height})$.

[EXERCISE : how to implement RANK ?

Tim Roughgarden



Design and Analysis
of Algorithms I

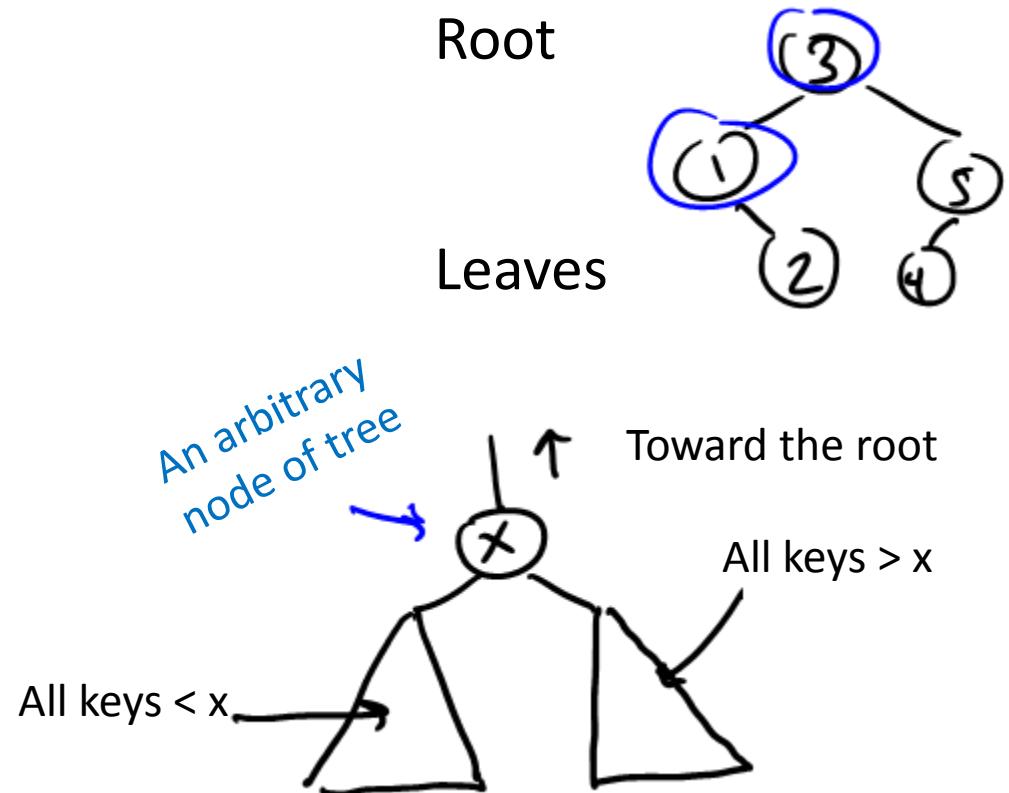
Data Structures

Red-Black Trees

Binary Search Tree Structure

- exactly one node per key
- most basic version :
 - each node has
 - left child pointer
 - right child pointer
 - parent pointer

SEARCH TREE PROPERTY :
(should hold at every node of the search tree)



The Height of a BST

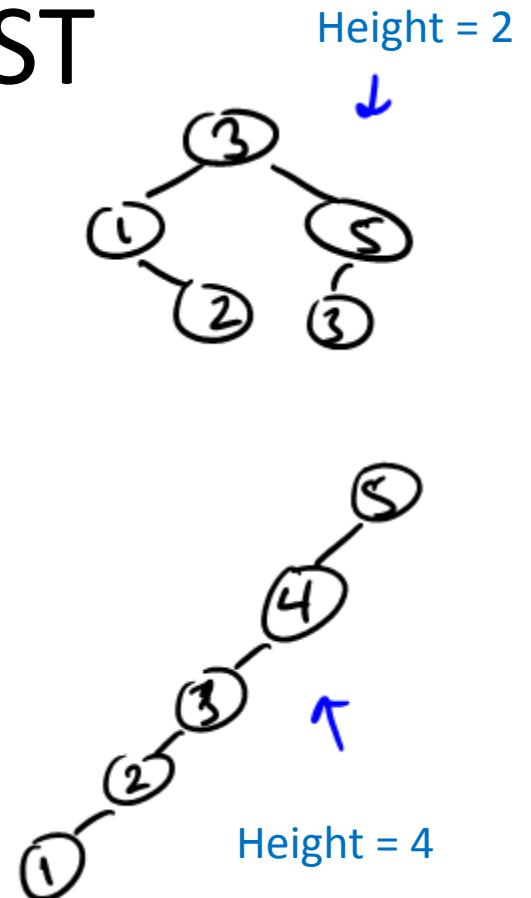
Note : many possible trees for a set of keys.

Note : height could be anywhere from $\sim \log_2 n$ to $\sim n$

(aka depth) longest root-leaf path

Worst case, a chain

Best case, perfectly balanced



Balanced Search Trees

Idea : ensure that height is always $O(\log(n))$ [best possible]

⇒ Search / Insert / Delete / Min / Max / Pred / Succ will then run in $O(\log(n))$ time [$n = \#$ of keys in tree]

Example : red-black trees [Bayes '72, Guibas-Sedgewick '78]

[see also AUL trees, splay trees, B trees]

Red-Black Invariants

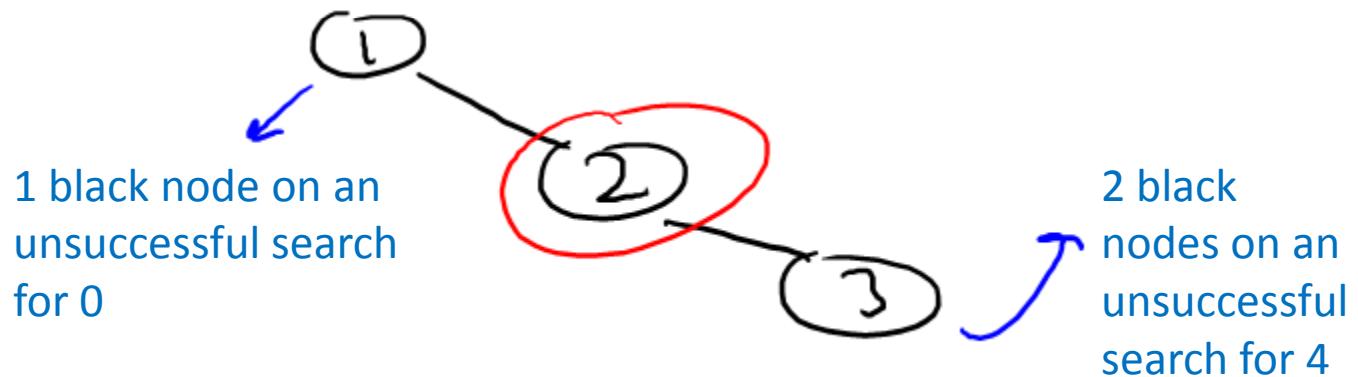
1. Each node red or black
2. Root is black
3. No 2 reds in a row
[red node => only black children]
4. Every root-NULL path has same number of black nodes

Like in an
unsuccessful search

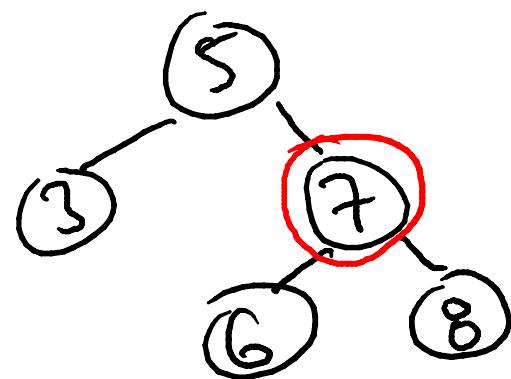
Example #1

Claim : a chain of length 3 cannot be a red-black tree

Proof



Example #2

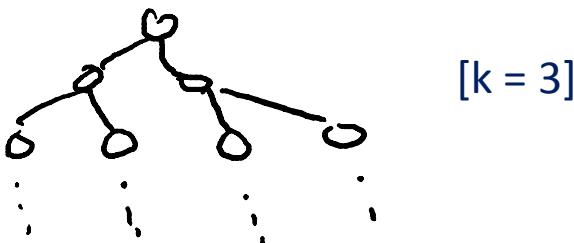


Height Guarantee

Claim : every red-black tree with n nodes has height $\leq 2 \log_2(n + 1)$

Proof : Observation : if every root-NUL path has $\geq k$ nodes, then tree includes (at the top) a perfectly balanced search tree of depth $k-1$.

=> Size n of the tree
must Be at least $2^k - 1$



Height Guarantee (con'd)

Story so far : size $n \geq 2^k - 1$, where $k = \text{minimum } \# \text{ of nodes on root - NULL path}$

$$\Rightarrow k \leq \log_2(n + 1)$$

Thus : in a red-black tree with n nodes, there is a root-NULL path with at most $\log_2(n + 1)$ black nodes.

By 4th Invariant : every root-NULL path has $\leq \log_2(n + 1)$ black nodes

By 3rd Invariant : every root-NULL path has $\leq 2 \log_2(n + 1)$ total nodes.

Q.E.D.

Which of the search tree operations have to be re-implemented so that the Red-Black invariants are maintained?

- Search
- Delete
- Insert and Delete
- None of the above



Data Structures

Rotations

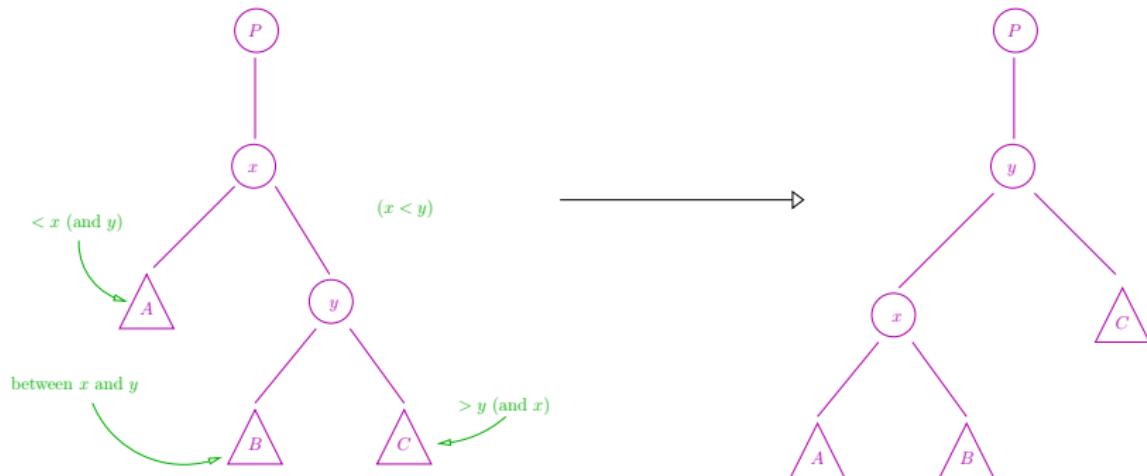
Design and Analysis
of Algorithms I

Left Rotations

Key primitive: *Rotations.* (Common to all balanced search tree implementations: red-black, AVL, B+, etc.)

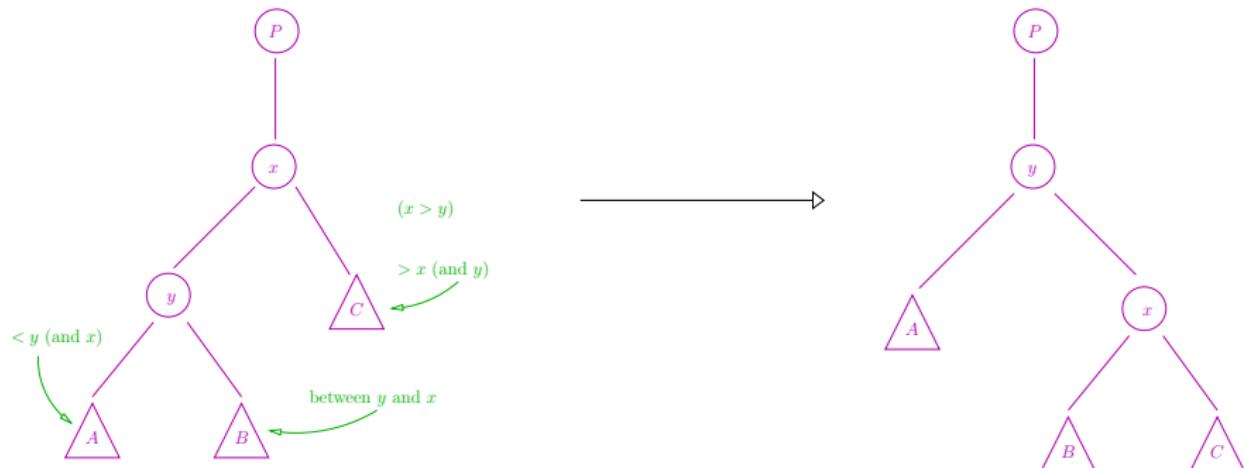
Idea: Locally rebalance subtrees at a node in $O(1)$ time.

Left rotation: (of a parent x and right child y)



Right Rotations

Right rotation:



Nice properties: Search tree property maintain, can implement in $O(1)$ time.



Data Structures

Insertion In A Red-Black Tree

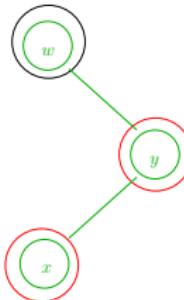
Design and Analysis
of Algorithms I

High-Level Plan

Idea for Insert/Delete: Proceed as in a normal binary search tree, then recolor and/or perform rotations until invariants are restored.

Insert(x):

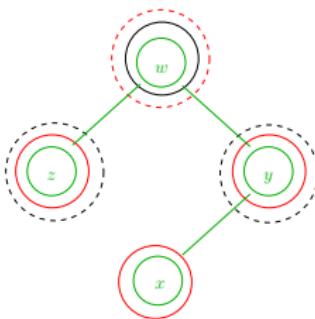
1. Insert x as usual (makes x a leaf).
2. Try coloring x red.
3. If x 's parent y is black, done.
4. Else y is red $\Rightarrow y$ has a black parent w .



Insertion

Case 1

- Case 1:** The other child z of x 's grandparent w is also red.
⇒ Recolor y, z black and w red. [key point: does not break invariant (4)]
⇒ Either restores invariant (3) or propagates the double red upward.
⇒ Can only happen $O(\log n)$ times. [If you reach the root, recolor it black ⇒ Preserves invariant (4)].



Case 2

Case 2: Let x, y be the current double-red, x the deeper node. Let $w = x$'s grandparent. Suppose w 's other child is NULL or is a black node z .

Exercise/case analysis (details omitted): Can eliminate double-red [\Rightarrow All invariants satisfied] in $O(1)$ time via 2-3 rotations+ recolorings.

