



Minimum Spanning Trees

Algorithms: Design
and Analysis, Part II

Problem Definition

Overview

Informal Goal: Connect a bunch of points together as cheaply as possible.

Applications: Clustering (more later), networking.

Blazingly Fast Greedy Algorithms:

- Prim's Algorithm [1957; also Dijkstra 1959, Jarnik 1930]
- Kruskal's algorithm [1956]

⇒ $O(m \log n)$ time (using suitable data structures)



Problem Definition

vertices

edges


Input: Undirected graph $G = (V, E)$ and a cost c_e for each edge $e \in E$.

- Assume adjacency list representation (see Part I for details)
- OK if edge costs are negative


Output: minimum cost tree $T \subseteq E$ that spans all vertices.

i.e., sum of edge costs

i.e.: (1) T has no cycles, (2) the subgraph (V, T) is connected (i.e., contains path between each pair of vertices).



(disallowed)



Standing Assumptions

Assumption #1: Input graph G is connected.

- Else no spanning trees.
- Easy to check in preprocessing (e.g., depth-first search).

Assumption #2: Edge costs are distinct.

- Prim + Kruskal remain correct with ties (which can be broken arbitrarily).
- Correctness proof a bit more annoying (will skip).



Algorithms: Design
and Analysis, Part II


Minimum Spanning Trees

Prim's MST Algorithm

Example

[Purple edges = minimum spanning tree]

(Compare to Dijkstra's shortest-path algorithm)



Prim's MST Algorithm

- Initialize $X = \{s\}$ [$s \in V$ chosen arbitrarily]
- $T = \emptyset$ [invariant: X = vertices spanned by tree-so-far T]
- While $X \neq V$
 - Let $e = (u, v)$ be the cheapest edge of G with $u \in X$, $v \notin X$.
 - Add e to T
 - Add v to X .

While loop: Increase # of spanned vertices in cheapest way possible.

Correctness of Prim's Algorithm

Theorem: Prim's algorithm always computes an MST.

Part I: Computes a spanning tree T^* .

[Will use basic properties of graphs and spanning trees] (Useful also in Kruskal's MST algorithm)

Part II: T^* is an MST.

[Will use the "Cut Property"] (Useful also in Kruskal's MST algorithm)

Later: Fast [$O(m \log n)$] implementation using heaps.



Algorithms: Design
and Analysis, Part II


Minimum Spanning Trees

Correctness of Prim's
Algorithm (Part I)

Cuts

Claim: Prim's algorithm outputs a spanning tree.

Definition: A cut of a graph $G = (V, E)$ is a partition of V into 2 non-empty sets.



Quiz on Cuts


Question: Roughly how many cuts does a graph with n vertices have?

- A) n C) 2^n (for each vertex, choose whether in A or in B)
B) n^2 D) n^n


Empty Cut Lemma

Empty Cut Lemma: A graph is not connected $\iff \exists$ cut (A, B) with no crossing edges.

Proof: (\Leftarrow) Assume the RHS. Pick any $u \in A$ and $v \in B$. Since no edges cross (A, B) there is no u, v path in G . $\Rightarrow G$ not connected.




(\Rightarrow) Assume the LHS. Suppose G has no $u - v$ path. Define
 $A = \{\text{Vertices reachable from } u \text{ in } G\}$ (u 's connected component)
 $B = \{\text{All other vertices}\}$ (all other connected components)
Note: No edges cross cut (A, B) (otherwise A would be bigger!)



Two Easy Facts

Double-Crossing Lemma: Suppose the cycle $C \subseteq E$ has an edge crossing the cut (A, B) : then so does some other edge of C .



Lonely Cut Corollary: If e is the only edge crossing some cut (A, B) , then it is not in any cycle. [If it were in a cycle, some other edge would have to cross the cut!]

Proof of Part I

Claim: Prim's algorithm outputs a spanning tree.

[Not claiming MST yet]

Proof: (1) Algorithm maintains invariant that T spans X
[straightforward induction - you check]



(2) Can't get stuck with $X \neq V$

[otherwise the cut $(X, V - X)$ must be empty; by Empty Cut Lemma input graph G is disconnected]

(3) No cycles ever get created in T . Why? Consider any iteration, with current sets X and T . Suppose e gets added.

Key point: e is the first edge crossing $(X, V - X)$ that gets added to $T \Rightarrow$ its addition can't create a cycle in T (by Lonely Cut Corollary). **QED!**



Algorithms: Design
and Analysis, Part II

Minimum Spanning Trees

Correctness of Prim's
Algorithm (Part II)

Correctness of Prim's Algorithm


Theorem: Prim's algorithm always outputs a minimum-cost spanning tree.

Key Question: When is it “safe” to include an edge in the tree-so-far?

The Cut Property

CUT PROPERTY: Consider an edge e of G . Suppose there is a cut (A, B) such that e is the cheapest edge of G that crosses it. Then e belongs to **the** MST of G .

Turns out MST is unique if edge costs are distinct



Cut Property Implies Correctness

Claim: Cut Property \Rightarrow Prim's algorithm is correct.

Proof: By previous video, Prim's algorithm outputs a spanning tree T^* .

Key point: Every edge $e \in T^*$ is explicitly justified by the Cut Property.

$\Rightarrow T^*$ is a subset of the MST

\Rightarrow Since T^* is already a spanning tree, it must be the MST

QED!



Minimum Spanning Trees

Algorithms: Design
and Analysis, Part II

Proof of the Cut
Property

The Cut Property

Assumption: Distinct edge costs.

CUT PROPERTY: Consider an edge e of G . Suppose there is a cut (A, B) such that e is the cheapest edge of G that crosses it. Then e belongs to the MST of G .

Proof Plan

Will argue by contradiction, using an exchange argument.


[Compare to scheduling application]

Suppose there is an edge e that is the cheapest one crossing a cut (A, B) , yet e is not in the MST T^* .

Idea: Exchange e with another edge in T^* to make it even cheaper (contradiction).

Question: Which edge to exchange e with?

Attempted Exchange




Note: Since T^* is connected, must construct an edge $f (\neq e)$ crossing (A, B) .

Idea: Exchange e and f to get a spanning tree cheaper than T^* (contradiction).

Exchanging Edges

Question: Let T^* be a spanning tree of G , $e \notin T^*$, $f \in T^*$. Is $T^* \cup \{e\} - \{f\}$ a spanning tree of G ?

- A) Yes always
- B) No never
- C) If e is the cheapest edge crossing some cut, then yes
- D) Maybe, maybe not (depending on the choice of e and f)



Exchange e, f :



(not a spanning tree)

Exchange e, e' :




(a spanning tree)

Smart Exchanges

Hope: Can always find suitable edge e' so that exchange yields bona fide spanning tree of G .

How? Let C = cycle created by adding e to T^* .



By the Double-Crossing Lemma: Some other edge e' of C [with $e' \neq e$ and $e' \in T^*$] crosses (A, B) .

You check: $T = T^* \cup \{e\} - \{e'\}$ is also a spanning tree.

Since $c_e < c_{e'}$, T cheaper than purported MST T^* , contradiction.



Minimum Spanning Trees

Algorithms: Design
and Analysis, Part II

Fast Implementation
of Prim's Algorithm

Running Time of Prim's Algorithm

- Initialize $X = \{s\}$ [$s \in V$ chosen arbitrarily]
- $T = \emptyset$ [invariant: X = vertices spanned by tree-so-far T]
- While $X \neq V$
 - Let $e = (u, v)$ be the cheapest edge of G with $u \in X, v \notin X$.
 - Add e to T , add v to X .

Running time of straightforward implementation:

- $O(n)$ iterations [where $n = \#$ of vertices]
 - $O(m)$ time per iteration [where $m = \#$ of edges]
- $\Rightarrow O(mn)$ time


BUT CAN WE DO BETTER?

Prim's Algorithm with Heaps

[Compare to fast implementation of Dijkstra's algorithm]

Invariant #1: Elements in heap = vertices of $V - X$.

Invariant #2: For $v \in V - X$, $\text{key}[v]$ = cheapest edge (u, v) with $u \in X$ (or $+\infty$ if no such edges exist).




Check: Can initialize heap with $O(m + n \log n) = O(m \log n)$ preprocessing.

To compare keys $n - 1$ Inserts $m \geq n - 1$ since G connected

Note: Given invariants, Extract-Min yields next vertex $v \notin X$ and edge (u, v) crossing $(X, V - X)$ to add to X and T , respectively. \diamond

Quiz: Issue with Invariant #2

Question: What is: (i) current value of $\text{key}[v]$ (ii) current value of $\text{key}[w]$ (iii) value of $\text{key}[w]$ after one more iteration of Prim's algorithm?




A) 11, 10, 4 C) 2, 10, 1

B) 2, 10, 10 D) 2, 10, 2

Maintaining Invariant #2

Issue: Might need to recompute some keys to maintain Invariant #2 after each Extract-Min.



Pseudocode: When v added to X :

- For each edge $(v, w) \in E$:
 - If $w \in V - X \rightarrow$ The only whose key might have changed (Update key if needed):
 - Delete w from heap
 - Recompute $\text{key}[w] := \min\{\text{key}[w], c_{vw}\}$
 - Re-Insert into heap

Subtle point/exercise:

Think through book-keeping needed to pull this off



Running Time with Heaps

- Dominated by time required for heap operations
 - $(n - 1)$ Inserts during preprocessing
 - $(n - 1)$ Extract-Mins (one per iteration of while loop)
 - Each edge (v, w) triggers one Delete/Insert combo
[When its first endpoint is sucked into X]
- $\Rightarrow O(m)$ heap operations [Recall $m \geq n - 1$ since G connected]
- $\Rightarrow O(m \log n)$ time [As fast as sorting!]



Minimum Spanning Trees

Algorithms: Design
and Analysis, Part II

Fast Implementation
of Prim's Algorithm

Running Time of Prim's Algorithm

- Initialize $X = \{s\}$ [$s \in V$ chosen arbitrarily]
- $T = \emptyset$ [invariant: X = vertices spanned by tree-so-far T]
- While $X \neq V$
 - Let $e = (u, v)$ be the cheapest edge of G with $u \in X, v \notin X$.
 - Add e to T , add v to X .

Running time of straightforward implementation:

- $O(n)$ iterations [where $n = \#$ of vertices]
 - $O(m)$ time per iteration [where $m = \#$ of edges]
- $\Rightarrow O(mn)$ time


BUT CAN WE DO BETTER?

Prim's Algorithm with Heaps

[Compare to fast implementation of Dijkstra's algorithm]

Invariant #1: Elements in heap = vertices of $V - X$.

Invariant #2: For $v \in V - X$, $\text{key}[v]$ = cheapest edge (u, v) with $u \in X$ (or $+\infty$ if no such edges exist).




Check: Can initialize heap with $O(m + n \log n) = O(m \log n)$ preprocessing.

To compare keys $n - 1$ Inserts $m \geq n - 1$ since G connected

Note: Given invariants, Extract-Min yields next vertex $v \notin X$ and edge (u, v) crossing $(X, V - X)$ to add to X and T , respectively. \diamond

Quiz: Issue with Invariant #2

Question: What is: (i) current value of $\text{key}[v]$ (ii) current value of $\text{key}[w]$ (iii) value of $\text{key}[w]$ after one more iteration of Prim's algorithm?




A) 11, 10, 4 C) 2, 10, 1

B) 2, 10, 10 D) 2, 10, 2

Maintaining Invariant #2

Issue: Might need to recompute some keys to maintain Invariant #2 after each Extract-Min.



Pseudocode: When v added to X :

- For each edge $(v, w) \in E$:
 - If $w \in V - X \rightarrow$ The only whose key might have changed (Update key if needed):
 - Delete w from heap
 - Recompute $\text{key}[w] := \min\{\text{key}[w], c_{vw}\}$
 - Re-Insert into heap

Subtle point/exercise:

Think through book-keeping needed to pull this off



Running Time with Heaps

- Dominated by time required for heap operations
 - $(n - 1)$ Inserts during preprocessing
 - $(n - 1)$ Extract-Mins (one per iteration of while loop)
 - Each edge (v, w) triggers one Delete/Insert combo
[When its first endpoint is sucked into X]
- $\Rightarrow O(m)$ heap operations [Recall $m \geq n - 1$ since G connected]
- $\Rightarrow O(m \log n)$ time [As fast as sorting!]



Algorithms: Design
and Analysis, Part II

Minimum Spanning Trees

Kruskal's MST
Algorithm

MST Review


Input: Undirected graph $G = (V, E)$, edge costs c_e .

Output: Min-cost spanning tree (no cycles, connected).

Assumptions: G is connected, distinct edge costs.

Cut Property: If e is the cheapest edge crossing some cut (A, B) , then e belongs to the MST.

Example



Kruskal's MST Algorithm

- Sort edges in order of increasing cost
[Rename edges $1, 2, \dots, m$ so that $c_1 < c_2 < \dots < c_m$]
- $T = \emptyset$
- For $i = 1$ to m
 - If $T \cup \{i\}$ has no cycles
 - Add i to T
- Return T



Algorithms: Design
and Analysis, Part II

Minimum Spanning Trees

Correctness of
Kruskal's Algorithm

Correctness of Kruskal (Part I)

Theorem: Kruskal's algorithm is correct.

Proof: Let T^* = output of Kruskal's algorithm on input graph G .


(1) Clearly T^* has no cycles.

(2) T^* is connected. Why?

(2a) By Empty Cut Lemma, only need to show that T^* crosses every cut.

(2b) Fix a cut (A, B) . Since G connected at least one of its edges crosses (A, B) .

Key point: Kruskal will include first edge crossing (A, B) that it sees [by Lonely Cut Corollary, cannot create a cycle]



Correctness of Kruskal (Part II)

(3) Every edge of T^* satisfied by the Cut Property. (Implies T^* is the MST)


Reason for (3): Consider iteration where edge (u, v) added to current set T . Since $T \cup \{(u, v)\}$ has no cycle, T has no $u - v$ path.

$\Rightarrow \exists$ empty cut (A, B) separating u and v . (As in proof of Empty Cut Lemma)

\Rightarrow By (2b), no edges crossing (A, B) were previously considered by Kruskal's algorithm.

$\Rightarrow (u, v)$ is the first (+ hence the cheapest!) edge crossing (A, B) .

$\Rightarrow (u, v)$ justified by the Cut Property. QED





Algorithms: Design
and Analysis, Part II

Minimum Spanning Trees

Implementing
Kruskal's Algorithm
via Union-Find

Kruskal's MST Algorithm

- Sort edges in order of increasing cost. ($O(m \log n)$, recall $m = O(n^2)$ assuming nonparallel edges)
- $T = \emptyset$
 - For $i = 1$ to m ($O(m)$ iterations)
 - If $T \cup \{i\}$ has no cycles ($O(n)$ time to check for cycle [Use BFS or DFS in the graph (V, T) which contains $\leq n - 1$ edges])
 - Add i to T
- Return T

Running time of straightforward implementation: ($m = \#$ of edges, $n = \#$ of vertices) $O(m \log n) + O(mn) = O(mn)$



Plan: Data structure for $O(1)$ -time cycle checks $\Rightarrow O(m \log n)$ time.

The Union-Find Data Structure

Raison d'être of union-find data structure: Maintain partition of a set of objects.

FIND(X): Return name of group that X belongs to.

UNION(C_i, C_j): Fuse groups C_i, C_j into a single one.



Why useful for Kruskal's algorithm: Objects = vertices


- Groups = Connected components w.r.t. chosen edges T .
- Adding new edge (u, v) to $T \iff$ Fusing connected components of u, v .

Union-Find Basics

Motivation: $O(1)$ -time cycle checks in Kruskal's algorithm.

Idea #1: - Maintain one linked structure per connected component of (V, T) .

- Each component has an arbitrary leader vertex.



Invariant: Each vertex points to the leader of its component [“name” of a component inherited from leader vertex]

Key point: Given edge (u, v) , can check if u & v already in same component in $O(1)$ time. [if and only if leader pointers of u, v match, i.e., $\text{FIND}(u) = \text{FIND}(v)$] $\Rightarrow O(1)$ -time cycle checks!

Maintaining the Invariant

Note: When new edge (u, v) added to T , connected components of u & v merge.

Question: How many leader pointer updates are needed to restore the invariant in the worst case?

- A) $\Theta(1)$
- B) $\Theta(\log n)$
- C) $\Theta(n)$ (e.g., when merging two components with $n/2$ vertices each)
- D) $\Theta(m)$

Maintaining the Invariant (con'd)

Idea #2: When two components merge, have smaller one inherit the leader of the larger one. [Easy to maintain a size field in each component to facilitate this]

Question: How many leader pointer updates are now required to restore the invariant in the worst case?

- A) $\Theta(1)$
- B) $\Theta(\log n)$
- C) $\Theta(n)$ (for same reason as before, i.e., when merging two components with $n/2$ vertices each)
- D) $\Theta(m)$

Updating Leader Pointers

But: How many times does a single vertex v have its leader pointer updated over the course of Kruskal's algorithm?

- A) $\Theta(1)$
- B) $\Theta(\log n)$
- C) $\Theta(n)$
- D) $\Theta(m)$

Reason: Every time v 's leader pointer gets updated, population of its component at least doubles \Rightarrow Can only happen $\leq \log_2 n$ times.

Running Time of Fast Implementation

Scorecard:

$O(m \log n)$ time for sorting

$O(m)$ times for cycle checks [$O(1)$ per iteration]

$O(n \log n)$ time overall for leader pointer updates

$O(m \log n)$ total (Matching Prim's algorithm)



Minimum Spanning Trees

Algorithms: Design
and Analysis, Part II

State-of-the-Art and
Open Questions


State-of-the-Art MST Algorithms

Question: Can we do better than $O(m \log n)$? (Running time of Prim/Kruskal.)

Answer: Yes!

$O(m)$ randomized algorithm [Karger-Klein-Tarjan JACM 1995]

$O(m \alpha(n))$ deterministic [Chazelle JACM 2000]



“Inverse Ackerman Function”: In particular, grows much slower than $\log^* n := \#$ of times you can apply \log to n until result drops below 1 (inverse of “tower function” $2^{2^{\dots^2}}$)

Open Questions

Weirdest of all: [Pettie/Ramachandran JACM 2002] Optimal deterministic MST algorithm, but precise asymptotic running time is unknown! [Between $\Theta(m)$ and $\Theta(m\alpha(n))$, but don't know where]

Open Questions:

- Simple randomized $O(m)$ -time algorithm for MST [Sufficient: Do this just for the “MST verification” problem]
- Is there a deterministic $O(m)$ -time algorithm?

Further reading: [Eisner 97]