




Advanced Union-Find

Algorithms: Design
and Analysis, Part II

Lazy Unions

The Union-Find Data Structure

Raison d'être: Maintain a partition of a set X .




FIND: Given $x \in X$, return name of x 's group.

UNION: Given x & y , merge groups containing them.

Previous solution (for Kruskal's MST algorithm)

- Each $x \in X$ points directly to the "leader" of its group.




- $O(1)$ FIND [just return x 's leader]
- $O(n \log n)$ total work for n UNIONS [when 2 groups merge, smaller group inherits leader of larger one]


Lazy Unions

New idea: Update only one pointer each merge!

How?




In array representation: (Where $A[i] \leftrightarrow$ name of i 's parent)



How to Merge?

In general: When two groups merge in a UNION, make one group's leader (i.e., root of the tree) a child of the other one.



Pro: UNION reduces to 2 FINDS [$r_1 = \text{FIND}(x)$, $r_2 = \text{FIND}(y)$] and $O(1)$ extra work [link r_1, r_2 together]

Con: To recover leader of an object, need to follow a path of parent pointers [not just one!]

⇒ Not clear if FIND still takes $O(1)$ time.




Algorithms: Design
and Analysis, Part II

Advanced Union-Find

Union by Rank

The Lazy Union Implementation

New implementation: Each object $x \in X$ has a parent field.



Invariant: Parent pointers induce a collection of directed trees on X . (x is a root \iff $\text{parent}[x]=x$)

Initially: For all x , $\text{parent}[x]=x$



FIND(x): Traverse parent pointers from x until you hit the root.


UNION(x, y): $s_1 = \text{FIND}(x)$; $s_2 = \text{FIND}(y)$; Reset parent of one of s_1, s_2 to be the other.

Quiz on Lazy Unions

Question: Suppose, in the UNION operation, we choose the new root arbitrarily from the two old ones. What is the worst-case running time of the FIND and UNION operations, respectively?

- A) $\Theta(1), \Theta(1)$
- B) $\Theta(\log n), \Theta(1)$
- C) $\Theta(\log n), \Theta(\log n)$
- D) $\Theta(n), \Theta(n)$


Issue: Scraggly trees:



Union by Rank

Ranks: For each $x \in X$, maintain field $\text{rank}[x]$.

[In general $\text{rank}[x] = 1 + (\text{max rank of } x\text{'s children})$]



Invariant (for now): For all $x \in X$, $\text{rank}[x] = \text{maximum number of hops from some leaf to } x$.

[Initially, $\text{rank}[x] = 0$ for all $x \in X$]


To avoid scraggly trees ("Union by Rank"): Given x & y :

- $s_1 = \text{FIND}(x)$, $s_2 = \text{FIND}(y)$
- If $\text{rank}[s_1] > \text{rank}[s_2]$ then set $\text{parent}[s_2]$ to s_1 else set $\text{parent}[s_1]$ to s_2 .

To-do: Update ranks to restore Invariant.

Quiz on Rank Updates

Question: Recall $s_1 = \text{FIND}(x)$, $s_2 = \text{FIND}(y)$. How do the ranks of s_1 & s_2 change after $\text{UNION}(x, y)$?



- A) Unchanged
- B) The one with larger rank goes up by 1
- C) The one with smaller rank goes up by 1
- D) No change unless ranks of s_1, s_2 were equal, in which case s_2 's rank goes up by 1



Algorithms: Design
and Analysis, Part II

Advanced Union-Find

Union by Rank -
Analysis

Properties of Ranks


Recall: Lazy Unions.

Invariant (for now): $\text{rank}[x] = \max \# \text{ of hops from a leaf to } x$.

[Note $\max_x \text{rank}[x] \approx \text{worst-case running time of FIND.}$]

Union by Rank: Make old root with smaller rank child of the root with the larger rank.

[Choose new root arbitrarily in case of a tie, and add 1 to its rank.]



Immediate from Invariant/Rank Maintenance:

- (1) For all objects x , $\text{rank}[x]$ only goes up over time
- (2) Only ranks of roots can go up
[once x a non-root, $\text{rank}[x]$ frozen for evermore]
- (3) Ranks strictly increase along a path to the root

Rank Lemma


Rank Lemma: Consider an arbitrary sequence of UNION (+FIND) operations. For every $r \in \{0, 1, 2, \dots\}$, there are at most $n/2^r$ objects with rank r .

Corollary: Max rank always $\leq \log_2 n$

Corollary: Worst-case running time of FIND, UNION is $O(\log n)$.
[With Union by Rank.]

Proof of Rank Lemma

Claim 1: If x, y have the same rank r , then their subtrees (objects from which can reach x, y) are disjoint.



Claim 2: The subtree of a rank- r object has size $\geq 2^r$.
[Note Claim 1 + Claim 2 imply the Rank Lemma.]

Proof of Claim 1: Will show contrapositive. Suppose subtrees of x, y have object z in common $\Rightarrow \exists$ path $z \rightarrow x, z \rightarrow y$
 \Rightarrow One of x, y is an ancestor of the other
 \Rightarrow The ancestor has strictly larger rank. [By property (3)]

QED (Claim 1)


Proof of Claim 2

Rank $r \Rightarrow$ Subtree size $\geq 2^r$

Base case: Initially all ranks = 0, all subtree sizes = 1

Inductive step: Nothing to prove unless the rank of some object changes (subtree sizes only go up).

Interesting case: UNION(x, y), with $s_1 = \text{FIND}(x)$, $s_2 = \text{FIND}(y)$, and $\text{rank}[s_1] = \text{rank}[s_2] = r \Rightarrow s_2$'s new rank = $r + 1$
 $\Rightarrow s_2$'s new subtree size = s_2 's old subtree size + s_1 's old subtree size (each at least 2^r by the inductive hypothesis) $\geq 2^{r+1}$. QED!





Algorithms: Design
and Analysis, Part II


Advanced Union-Find

Path Compression

Path Compression

Idea: Why bother traversing a leaf-root path multiple times?

Path compression: After $\text{FIND}(x)$, install shortcuts (i.e., revise parent pointers) to x 's root all along the $x \rightarrow \text{root}$ path.



In array representation:




Con: Constant-factor overhead to FIND (from “multitasking”).

Pro: Speeds up subsequent FIND s. [But by how much?]

On Ranks

Important: Maintain all rank fields EXACTLY as without path compression.

- Ranks initially all 0
- In UNION, new root = old root with bigger rank
- When merging two nodes of common rank r , reset new root's rank to $r + 1$



Bad news: Now $\text{rank}[x]$ is only an upper bound on the maximum number of hops on a path from a leaf to x
(which could be much less)

Good news: Rank Lemma still holds ($\leq n/2^r$ objects with rank r)

Also: Still always have $\text{rank}[\text{parent}[x]] > \text{rank}[x]$ for all non-roots x

Hopcroft-Ullman Theorem

Theorem: [Hopcroft-Ullman 73] With Union by Rank and path compression, m Union+Find operations take $O(m \log^* n)$ time, where $\log^* n = \text{the number of times you need to apply log to } n \text{ before the result is } \leq 1$.

Quiz on \log^*

Question: What is $\log^*(2^{65536})$?

- A) 2
- B) 5
- C) 16
- D) 65536

In general: $\log^*(2^{2 \dots t \text{ times } \dots^2}) = t$

Measuring Progress



Algorithms: Design
and Analysis, Part II

Advanced Union-Find

Path Compression: The
Hopcroft-Ullman Analysis

Hopcroft-Ullman Theorem

Theorem: [Hopcroft-Ullman 73] With Union by Rank and path compression, m UNION+FIND operations take $O(m \log^* n)$ time, where $\log^* n$ = the number of times you need to apply log to n before the result is ≤ 1 .

[Will focus on interesting case where $m = \Omega(n)$]

Measuring Progress


Intuition: Installing shortcuts should significantly speed up subsequent FINDs+UNIONS.

Question: How to track this progress and quantify the benefit?

Idea: Consider a non-root object x . \longrightarrow **Recall:** $\text{rank}[x]$ frozen


Progress measure: $\text{rank}[\text{parent}[x]] - \text{rank}[x]$

Path compression increases this progress measure: If x has old parent p , new parent $p' \neq p$, then $\text{rank}[p'] > \text{rank}[p]$.



Proof Setup

Rank blocks: $\{0\}, \{1\}, \{2, 3, 4\}, \{5, \dots, 2^4\}, \{17, 18, \dots, 2^{16}\},$
 $\{65537, \dots, 2^{65536}\}, \dots, \{\dots, n\}$



The diagram shows two blue numbers, 16 and 65536, with arrows pointing to the corresponding exponents in the rank blocks sequence. The arrow from 16 points to the 2^4 term, and the arrow from 65536 points to the 2^{65536} term.

Note: There are $O(\log^* n)$ different rank blocks.

Semantics: Traversal $x \rightarrow \text{parent}(x)$ is “fast progress” \iff
 $\text{rank}[\text{parent}[x]]$ in larger block than $\text{rank}[x]$

Definition: At a given point in time, call object x good if

- (1) x or x 's parent is a root OR
- (2) $\text{rank}[\text{parent}[x]]$ in larger block than $\text{rank}[x]$

x is bad otherwise.


Proof of Hopcroft-Ullman

Point: Every FIND visits only $O(\log^* n)$ good nodes $[2 + \# \text{ of rank blocks} = O(\log^* n)]$

Upshot: Total work done during m operations $= O(m \log^* n)$
(visits to good objects) + total $\#$ of visits to bad nodes (need to bound globally by separate argument)

Consider: A rank block $\{k + 1, k + 2, \dots, 2^k\}$.

Note: When a bad node is visited



its parent is changed to one with strictly larger rank \Rightarrow Can only happen 2^k times before x becomes good (forevermore).

Proof of Hopcroft-Ullman II

Total work: $O(m \log^* n) + O(\text{\# visits to bad nodes})$.

$\leq n$ for each of $O(\log^* n)$ rank blocks

Consider: A rank block $\{k+1, k+2, \dots, 2^k\}$.

Last slide: For each object x with final rank in this block, $\text{\# visits to } x \text{ while } x \text{ is bad}$ is $\leq 2^k$.

Rank Lemma: Total number of objects x with final rank in this rank block is $\sum_{i=k+1}^{2^k} n/2^i \leq n/2^k$.

$\leq n$ visits to bad objects in this rank block.

Recall: Only $O(\log^* n)$ rank blocks.

Total work: $O((m+n) \log^* n)$.



Algorithms: Design
and Analysis, Part II

Advanced Union-Find

Path Compression: The
Hopcroft-Ullman Analysis

Hopcroft-Ullman Theorem

Theorem: [Hopcroft-Ullman 73] With Union by Rank and path compression, m UNION+FIND operations take $O(m \log^* n)$ time, where $\log^* n$ = the number of times you need to apply log to n before the result is ≤ 1 .

[Will focus on interesting case where $m = \Omega(n)$]

Measuring Progress


Intuition: Installing shortcuts should significantly speed up subsequent FINDs+UNIONS.

Question: How to track this progress and quantify the benefit?

Idea: Consider a non-root object x . \longrightarrow **Recall:** $\text{rank}[x]$ frozen


Progress measure: $\text{rank}[\text{parent}[x]] - \text{rank}[x]$

Path compression increases this progress measure: If x has old parent p , new parent $p' \neq p$, then $\text{rank}[p'] > \text{rank}[p]$.



Proof Setup

Rank blocks: $\{0\}, \{1\}, \{2, 3, 4\}, \{5, \dots, 2^4\}, \{17, 18, \dots, 2^{16}\},$
 $\{65537, \dots, 2^{65536}\}, \dots, \{\dots, n\}$



The diagram shows two blue numbers, 16 and 65536, with arrows pointing to the corresponding exponents in the rank blocks sequence. An arrow points from 16 to 2^4 in the block $\{5, \dots, 2^4\}$. Another arrow points from 65536 to 2^{65536} in the block $\{65537, \dots, 2^{65536}\}$.

Note: There are $O(\log^* n)$ different rank blocks.

Semantics: Traversal $x \rightarrow \text{parent}(x)$ is “fast progress” \iff
 $\text{rank}[\text{parent}[x]]$ in larger block than $\text{rank}[x]$

Definition: At a given point in time, call object x good if

- (1) x or x 's parent is a root OR
- (2) $\text{rank}[\text{parent}[x]]$ in larger block than $\text{rank}[x]$

x is bad otherwise.


Proof of Hopcroft-Ullman

Point: Every FIND visits only $O(\log^* n)$ good nodes $[2 + \# \text{ of rank blocks} = O(\log^* n)]$

Upshot: Total work done during m operations $= O(m \log^* n)$
(visits to good objects) + total $\#$ of visits to bad nodes (need to bound globally by separate argument)

Consider: A rank block $\{k + 1, k + 2, \dots, 2^k\}$.

Note: When a bad node is visited



its parent is changed to one with strictly larger rank \Rightarrow Can only happen 2^k times before x becomes good (forevermore).

Proof of Hopcroft-Ullman II

Total work: $O(m \log^* n) + O(\text{\# visits to bad nodes})$.

$\leq n$ for each of $O(\log^* n)$ rank blocks

Consider: A rank block $\{k+1, k+2, \dots, 2^k\}$.

Last slide: For each object x with final rank in this block, # visits to x while x is bad is $\leq 2^k$.

Rank Lemma: Total number of objects x with final rank in this rank block is $\sum_{i=k+1}^{2^k} n/2^i \leq n/2^k$.

$\leq n$ visits to bad objects in this rank block.

Recall: Only $O(\log^* n)$ rank blocks.

Total work: $O((m+n) \log^* n)$.



Algorithms: Design
and Analysis, Part II

Advanced Union-Find

The Ackermann Function

Tarjan's Bound

Theorem: [Tarjan 75] With Union by Rank and path compression, m UNION+FIND operations take $O(m\alpha(n))$ time, where $\alpha(n)$ is the inverse Ackerman function (will define in this video)

Proof in next video.

The Ackermann Function

Aside: Many different definitions, all more or less equivalent.

Will define $A_k(r)$ for all integers k and $r \geq 1$. (recursively)

Base case: $A_0(r) = r + 1$ for all $r \geq 1$.

In general: For $k, r \geq 1$:

$$\begin{aligned} A_k(r) &= \text{Apply } A_{k-1} \text{ } r \text{ times to } r \\ &= (A_{k-1} \circ A_{k-1} \circ \dots \circ A_{k-1})(r) \end{aligned}$$

r -fold composition



Quiz: A_1

Quiz: $A_1(r)$ corresponds to what function of r ?

- A) Successor ($r \mapsto r + 1$)
- B) Doubling ($r \mapsto 2r$)
- C) Exponentiation ($r \mapsto 2^r$)
- D) Tower function ($r \mapsto 2^{2^{\dots r \text{ times } \dots^2}}$)

$$A_1(r) = (A_0 \circ A_0 \circ \dots \circ A_0)(r) = 2r$$

(r -fold composition, add 1 each time)

Quiz: A_2

Quiz: What function does $A_2(r)$ correspond to?

A) $r \mapsto 4r$

B) $r \mapsto 2^r$

B) $r \mapsto r2^r$

D) $r \mapsto 2^{2 \dots r \text{ times} \dots 2}$

$A_2(r) = (A_1 \circ A_1 \circ \dots \circ A_1)(r) = r2^r$
(r -fold composition, doubles each time)

Quiz: A_3

Quiz: What is $A_3(2)$? Recall $A_2(r) = r2^r$

- A) 8
- B) 1024
- B) 2048
- D) Bigger than 2048

$$A_3(2) = A_2(A_2(2)) = A_2(8) = 82^8 = 2^{11} = 2048$$

In general: $A_3(r) = (A_2 \circ A_2 \circ \dots (r \text{ times}) \dots \circ A_2)(r) \geq$ a tower
of r 2's $= 2^{2^{\dots r \text{ times} \dots 2}}$

$$A_4$$

$$A_4(2) = A_3(A_3(2)) = A_3(2048) \geq 2^{2^{\dots \text{height } 2048} \dots^2}$$

In general: $A_4(r) = (A_3 \circ \dots \text{ } r \text{ times } \dots \circ A_3)(r) \approx$ iterated tower function (aka “wowzer” function)

The Inverse Ackermann Function

Definition: For every $n \geq 4$, $\alpha(n)$ = minimum value of k such that $A_k(2) \geq n$.

$$\alpha(n) = 1, n = 4 \text{ (} A_1(2) = 4 \text{)}$$

$$\alpha(n) = 2, n = 5, \dots, 8 \text{ (} A_2(2) = 8 \text{)}$$

$$\alpha(n) = 3, n = 9, 10, \dots, 2048$$

$$\alpha(n) = 4, n \text{ up to roughly a tower}$$

of 2's of height 2048

$$\alpha(n) = 5 \text{ for } n \text{ up to ???}$$

$$\log^* n = 1, n = \underline{2}$$

$$\log^* n = 2, n = 3, \underline{4}$$

$$\log^* n = 3, n = 5, \dots, \underline{16}$$

$$\log^* n = 4, n = 17, \underline{65536}$$

$$\log^* n = 5, n = 65537, \underline{2^{65536}}$$

$$\log^* n = 2048 \text{ for such } n$$



Algorithms: Design
and Analysis, Part II

Advanced Union-Find

Tarjan's Analysis

Tarjan's Bound

Theorem: [Tarjan 75] With Union by Rank and path compression, m UNION+FIND operations take $O(m\alpha(n))$ time, where $\alpha(n)$ is the inverse Ackerman function

Acknowledgement: Kozen, "Design and Analysis of Algorithms"

Building Blocks of Hopcroft-Ullman Analysis

Block #1: Rank Lemma (at most $n/2^r$ objects of rank r)

Block #2: Path compression \Rightarrow If x 's parent pointer updated from p to p' , then $\text{rank}(p') \geq \text{rank}(p) + 1$

New idea: Stronger version of building block #2. In most cases, rank of new parent much bigger than rank of old parent (not just by 1).

Quantifying Rank Gaps

Definition: Consider a non-root object x (so $\text{rank}[x]$ fixed forevermore)

Define $\delta(x) = \max$ value of k such that $\text{rank}[\text{parent}[x]] \geq A_k(\text{rank}[x])$

(Note $\delta(x)$ only goes up over time)

Examples: Always have $\delta(x) \geq 0$

$$\delta(x) \geq 1 \iff \text{rank}[\text{parent}[x]] \geq 2 \text{ rank}[x]$$

$$\delta(x) \geq 2 \iff \text{rank}[\text{parent}[x]] \geq \text{rank}[x] 2^{\text{rank}[x]}$$

Note: For all objects x with $\text{rank}[x] \geq 2$, then $\delta(x) \leq \alpha(n)$
[Since $A_{\alpha(n)}(2) \geq n$]

Good and Bad Objects

Definition: An object x is bad if all of the following hold:

- (1) x is not a root
- (2) $\text{parent}(x)$ is not a root
- (3) $\text{rank}(x) \geq 2$
- (4) x has an ancestor y with $\delta(y) = \delta(x)$

x is good otherwise.

Quiz

Question: What is the maximum number of good objects on an object-root path?


- A) $\Theta(1)$
- B) $\Theta(\alpha(n))$
- C) $\Theta(\log^* n)$
- D) $\Theta(\log n)$

\leq 1 root + 1 child of root
+ 1 object with rank 0
+ 1 object with rank 1
+ 1 object with $\delta(x) = k$
for each $k = 0, 1, 2, \dots, \alpha(n)$

Proof of Tarjan's Bound

Upshot: Total work of m operations = $O(m\alpha(n))$ (visits to good objects) + total # of visits to bad objects (will show is $O(n\alpha(n))$)

Main argument: Suppose a FIND operation visits a bad object x :



Path compression: x 's new parent will be p' or even higher.

$$\Rightarrow \text{rank}[x\text{'s new parent}] \geq \text{rank}[p'] \geq A_k(\text{rank}[y]) \geq A_k(\text{rank}[p])$$

ranks only go up

since $\delta(y) = k$

ranks only go up

Proof of Tarjan's Bound II

Point: Path compression (at least) applies the A_k function to $\text{rank}[x\text{'s parent}]$

Consequence: If $r = \text{rank}[x]$ (≥ 2), then after r such pointer updates we have

$$\text{rank}[x\text{'s parent}] \geq (A_k \circ \dots \text{ } r \text{ times } \dots \circ A_k)(r) = A_{k+1}(r)$$

Definition of Ackermann function

Thus: While x is bad, every r visits increases $\delta(x)$
 $\Rightarrow \leq r\alpha(n)$ visits to x while it's bad

Proof of Tarjan's Bound III

Recall: Total work of m operations is $O(m\alpha(n))$ (visits to good objects) + total # of visits to bad objects.

$$\leq \sum_{\text{objects } x} \text{rank}[x] \alpha(n)$$

$$= \alpha(n) \sum_{r \geq 0} r \cdot (\# \text{ of objects with rank } r)$$

$\leq n/2^r$ for each r by the Rank Lemma

$$= n\alpha(n) \sum_{r \geq 0} r/2^r \longrightarrow = O(1)$$

$$= O(n\alpha(n)). \quad \text{QED!}$$

Epilogue

“This is probably the first and maybe the only existing example of a simple algorithm with a very complicated running time. . . . I conjecture that there is no linear-time method, and that the algorithm considered here is optimal to within a constant factor.”

-Tarjan, “Efficiency of a Good But Non Linear Set Union Algorithm”, Journal of the ACM, 1975.

Conjecture proved by [Fredman/Saks 89]!



Algorithms: Design
and Analysis, Part II

Advanced Union-Find

Tarjan's Analysis

Tarjan's Bound

Theorem: [Tarjan 75] With Union by Rank and path compression, m UNION+FIND operations take $O(m\alpha(n))$ time, where $\alpha(n)$ is the inverse Ackerman function

Acknowledgement: Kozen, "Design and Analysis of Algorithms"

Building Blocks of Hopcroft-Ullman Analysis

Block #1: Rank Lemma (at most $n/2^r$ objects of rank r)

Block #2: Path compression \Rightarrow If x 's parent pointer updated from p to p' , then $\text{rank}(p') \geq \text{rank}(p) + 1$

New idea: Stronger version of building block #2. In most cases, rank of new parent much bigger than rank of old parent (not just by 1).

Quantifying Rank Gaps

Definition: Consider a non-root object x (so $\text{rank}[x]$ fixed forevermore)

Define $\delta(x) = \max$ value of k such that $\text{rank}[\text{parent}[x]] \geq A_k(\text{rank}[x])$

(Note $\delta(x)$ only goes up over time)

Examples: Always have $\delta(x) \geq 0$

$$\delta(x) \geq 1 \iff \text{rank}[\text{parent}[x]] \geq 2 \text{ rank}[x]$$

$$\delta(x) \geq 2 \iff \text{rank}[\text{parent}[x]] \geq \text{rank}[x] 2^{\text{rank}[x]}$$

Note: For all objects x with $\text{rank}[x] \geq 2$, then $\delta(x) \leq \alpha(n)$
[Since $A_{\alpha(n)}(2) \geq n$]

Good and Bad Objects

Definition: An object x is bad if all of the following hold:

- (1) x is not a root
- (2) $\text{parent}(x)$ is not a root
- (3) $\text{rank}(x) \geq 2$
- (4) x has an ancestor y with $\delta(y) = \delta(x)$

x is good otherwise.

Quiz

Question: What is the maximum number of good objects on an object-root path?


- A) $\Theta(1)$
- B) $\Theta(\alpha(n))$
- C) $\Theta(\log^* n)$
- D) $\Theta(\log n)$

\leq 1 root + 1 child of root
+ 1 object with rank 0
+ 1 object with rank 1
+ 1 object with $\delta(x) = k$
for each $k = 0, 1, 2, \dots, \alpha(n)$

Proof of Tarjan's Bound

Upshot: Total work of m operations = $O(m\alpha(n))$ (visits to good objects) + total # of visits to bad objects (will show is $O(n\alpha(n))$)

Main argument: Suppose a FIND operation visits a bad object x :



Path compression: x 's new parent will be p' or even higher.

$$\Rightarrow \text{rank}[x\text{'s new parent}] \geq \text{rank}[p'] \geq A_k(\text{rank}[y]) \geq A_k(\text{rank}[p])$$

ranks only go up

since $\delta(y) = k$

ranks only go up

Proof of Tarjan's Bound II

Point: Path compression (at least) applies the A_k function to $\text{rank}[x\text{'s parent}]$

Consequence: If $r = \text{rank}[x]$ (≥ 2), then after r such pointer updates we have

$$\text{rank}[x\text{'s parent}] \geq (A_k \circ \dots \text{ } r \text{ times } \dots \circ A_k)(r) = A_{k+1}(r)$$

Definition of Ackermann function

Thus: While x is bad, every r visits increases $\delta(x)$
 $\Rightarrow \leq r\alpha(n)$ visits to x while it's bad

Proof of Tarjan's Bound III

Recall: Total work of m operations is $O(m\alpha(n))$ (visits to good objects) + total # of visits to bad objects.

$$\leq \sum_{\text{objects } x} \text{rank}[x] \alpha(n)$$

$$= \alpha(n) \sum_{r \geq 0} r \cdot (\# \text{ of objects with rank } r)$$

$\leq n/2^r$ for each r by the Rank Lemma

$$= n\alpha(n) \sum_{r \geq 0} r/2^r \longrightarrow = O(1)$$

$$= O(n\alpha(n)). \quad \text{QED!}$$

Epilogue

“This is probably the first and maybe the only existing example of a simple algorithm with a very complicated running time. . . . I conjecture that there is no linear-time method, and that the algorithm considered here is optimal to within a constant factor.”

-Tarjan, “Efficiency of a Good But Non Linear Set Union Algorithm”, Journal of the ACM, 1975.

Conjecture proved by [Fredman/Saks 89]!