



Algorithms: Design
and Analysis, Part II

Huffman Codes

Introduction and
Motivation

Binary Codes

Binary code: Maps each character of an alphabet Σ to a binary string.

Example: $\Sigma = \text{a-z and various punctuation}$ (size 32 overall, say)

Obvious encoding: Use the 32 5-bit binary strings to encode this Σ (a fixed-length code)

Can we do better? Yes, if some characters of Σ are much more frequent than others, using a variable-length code.

Ambiguity

Example: Suppose $\Sigma = \{A,B,C,D\}$. Fixed-length encoding would be $\{00,01,10,11\}$.

Suppose instead we use the encoding $\{0,01,10,1\}$. What is 001 an encoding of?

- A) AB \rightarrow Leads to 001
- B) CD
- C) AAD \rightarrow Also leads to 001
- D) Not enough info to answer question

Prefix-Free Codes

Problem: With variable-length codes, not clear where one character ends + the next one begins.

Solution: Prefix-free codes - make sure that for every pair $i, j \in \Sigma$, neither of the encodings $f(i), f(j)$ is a prefix of the other.

Example: $\{0, 10, 110, 111\}$

Why useful? Can give shorter encodings with non-uniform character frequencies.

Example

Example:

| | | | |
|---|-----|----|-----|
| A | 60% | 00 | 0 |
| B | 25% | 01 | 10 |
| C | 10% | 10 | 110 |
| D | 5% | 11 | 111 |

| | | | |
|----------|-------------|--------------|----------------------------------|
| Σ | frequencies | fixed-length | variable-length (prefix free) |
|----------|-------------|--------------|----------------------------------|

Fixed-length encoding: 2 bits/character

Variable-length encoding: How many bits needed on average?

A) 1.5 B) 1.55 C) 2 D) 2.5

$$0.6 \cdot 1 + 0.25 \cdot 2 + (0.1 + 0.05) \cdot 3 = 1.55$$



Huffman Codes

Problem Definition


Algorithms: Design
and Analysis, Part II

Codes as Trees


Goal: Best binary prefix-free encoding for a given set of character frequencies.

Useful fact: Binary codes \leftrightarrow Binary trees


Examples: ($\Sigma = \{A,B,C,D\}$)



$\{00,01,10,11\}$



$\{0,01,10,1\}$




$\{0,10,110,111\}$

Prefix-Free Codes as Trees

- In general:**
- Left child edges \leftrightarrow "0", right child edges \leftrightarrow "1"
 - For each $i \in \Sigma$, exactly one node labeled "i"
 - Encoding of $i \in \Sigma \leftrightarrow$ Bits along path from node to the node "i"
 - Prefix-free \leftrightarrow Labelled nodes = the leaves
- [since prefixes \leftrightarrow one node an ancestor of another]

To decode: Repeatedly follow path from root until you hit a leaf.
[ex. 0110111 \mapsto ACD] (unambiguous since only leaves are labelled)



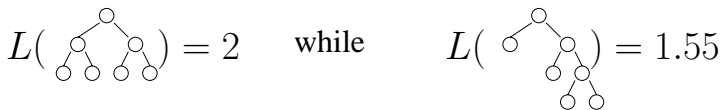
Note: Encoding length of $i \in \Sigma =$ depth of i in tree.

Problem Definition

Input: Probability p_i for each character $i \in \Sigma$.

Notation: If T = tree with leaves \leftrightarrow symbols of Σ , then **average encoding length** $L(T) = \sum_{i \in \Sigma} p_i \cdot [\text{depth of } i \text{ in } T]$

Example: If $p_A = 60\%$, $p_B = 25\%$, $p_C = 10\%$, $p_D = 5\%$, then

$$L(\text{balanced tree}) = 2 \quad \text{while} \quad L(\text{unbalanced tree}) = 1.55$$


Output: A binary tree T minimizing the average encoding length $L(\cdot)$.



Huffman Codes

A Greedy Algorithm


Algorithms: Design
and Analysis, Part II

Codes as Trees

Input: Probability p_i for each character $i \in \Sigma$.

Output: Binary tree (with leaves \leftrightarrow symbols of Σ) minimizing the average encoding length:

$$L(T) = \sum_{i \in \Sigma} p_i [\text{depth of } i \text{ in } T]$$

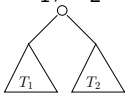


Building a Tree


Question: What's a principled approach for building a tree with leaves \leftrightarrow symbols of Σ ?

Natural but suboptimal idea: Top-down/divide+conquer.

- Partition Σ into Σ_1, Σ_2 each with $\approx 50\%$ of total frequency.
- Recursively compute T_1 for Σ_1 , T_2 for Σ_2 , return:



Huffman's (optimal) idea: Build tree bottom-up using successive mergers.




A Greedy Approach

Question: Which pair of symbols is “safe” to merge?

Observation: Final encoding length of $i \in \Sigma = \#$ of mergers its subtree endures.

[Each merger increases encoding length of participating symbols by 1]



Greedy heuristic: In first iteration, merge the two symbols with the smallest frequencies.

How to Recurse?

Suppose: 1st iteration of algorithm merges symbols a & b .

Idea: Replace symbols a, b by a new “meta-symbol” ab .

Question: What should be the frequency p_{ab} if this meta-symbol?


A) $\max\{p_a, p_b\}$

B) $\min\{p_a, p_b\}$

C) $p_a + p_b$ since ab is a proxy for “ a or b ” (intuitively)

D) $p_a - p_b$

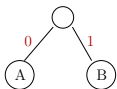
Example



Huffman's Algorithm

(Given frequencies p_i as input)

If $|\Sigma| = 2$ return




Let $a, b \in \Sigma$ have the smallest frequencies.

Let $\Sigma' = \Sigma$ with a, b replaced by new symbol ab .

Define $p_{ab} = p_a + p_b$.

Recursively compute T' (for the alphabet Σ')

Extend T' (with leaves $\leftrightarrow \Sigma'$) to a tree T with leaves $\leftrightarrow \Sigma$ by splitting leaf ab into two leaves a & b .



Return T



Algorithms: Design
and Analysis, Part II

Huffman Codes

Huffman's Algorithm: A
More Complex Example

Input and Steps 1 and 2

Input:

| | | | | | | |
|------------|---|---|---|---|---|---|
| Characters | A | B | C | D | E | F |
| Weights | 3 | 2 | 6 | 8 | 2 | 6 |

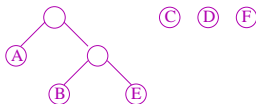
Step 1: Merge B and E:

| | | | | |
|---|----|---|---|---|
| A | BE | C | D | F |
| 3 | 4 | 6 | 8 | 6 |



Step 2: Merge A and BE:

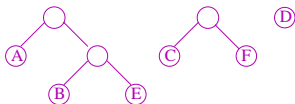
| | | | |
|-----|---|---|---|
| ABE | C | D | F |
| 7 | 6 | 8 | 6 |



Steps 3 and 4


Step 3: Merge C and F:

| | | |
|-----|----|---|
| ABE | CF | D |
| 7 | 12 | 8 |



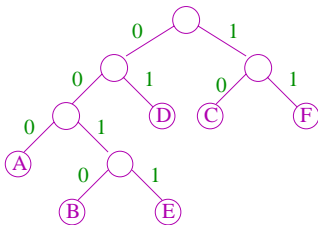
Step 2: Merge ABE and D:

| | |
|------|----|
| ABDE | CF |
| 15 | 12 |



Final Output

Final tree:



Corresponding code:

| | | | |
|---|------|---|----|
| A | 000 | D | 01 |
| B | 0010 | E | 10 |
| C | 0011 | F | 11 |



Huffman Codes

Correctness Proof

Algorithms: Design
and Analysis, Part II


Correctness of Huffman's Algorithm

Theorem: [Huffman 52] Huffman's algorithm computes a binary tree (with leaves \leftrightarrow symbols of Σ) that minimizes the average encoding length

$$L(T) = \sum_{i \in \Sigma} p_i [\text{depth of leaf } i \text{ in } T].$$

Proof: By induction on $n = |\Sigma|$. (Can assume $n \geq 2$.)

Base case: When $n = 2$, algorithm outputs the optimal tree.
(Needs 1 bit per symbol)



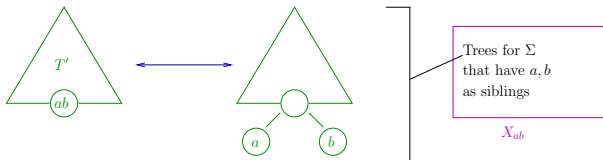
Inductive step: Fix input with $n = |\Sigma| > 2$.

By inductive hypothesis: Algorithm solves smaller subproblems (for Σ') optimally.

Inductive Step

Let $\Sigma' = \Sigma$ with a, b (symbols with smallest frequencies) replaced by meta-symbol ab . Define $p_{ab} = p_a + p_b$.

Recall: Exact correspondence between:



Important: For every such pair T' and T , $L(T) - L(T')$ is (after cancellation)

$$p_a [a's \text{ depth in } T] + p_b [b's \text{ depth in } T] - p_{ab} [ab's \text{ depth in } T'] =$$

Each is one more than

$$= p_a(d+1) + p_b(d+1) - (p_a + p_b)d = p_a + p_b, \text{ Independent of } T, T'!$$

Proof of Theorem

Inductive hypothesis: Huffman's algorithm computes a tree \hat{T}' that minimizes $L(T')$ for Σ' .

Upshot of last slide: Corresponding tree \hat{T} minimizes $L(T)$ for Σ over all trees in X_{ab} (i.e., where a & b are siblings)


Key lemma: [Completes proof of theorem] There is an optimal tree (for Σ) in X_{ab} . [i.e., a & b were “safe” to merge]

Intuition: Can make an optimal tree better by pushing a & b as deep as possible (since a, b have smallest frequencies).

Proof of Key Lemma

By exchange argument. Let T^* be any tree that minimizes $L(T)$ for Σ . Let x, y be siblings at the deepest level of T^* .

The exchange: Obtain \hat{T} from T^* by swapping $a \leftrightarrow x$, $b \leftrightarrow y$



Note: $\hat{T} \in X_{ab}$ (by choice of x, y).

To finish: Will show that $L(\hat{T}) \leq L(T^*)$

$\Rightarrow \hat{T}$ also optimal, completes proof]

Reason:

$$\begin{aligned}
 L(T^*) - L(\hat{T}) &= (p_x - p_a) \text{ [x's depth in } T^* - a\text{'s depth in } T^*] \\
 &+ (p_y - p_b) \text{ [y's depth in } T^* - b\text{'s depth in } T^*] \\
 &\geq 0 \quad \text{QED!}
 \end{aligned}$$

≥ 0 since a, b have smallest frequencies

≥ 0 by choice of x, y

Notes on Running Time

Naive implementation: $O(n^2)$ time, where $n = |\Sigma|$.

Speed ups: - Use a heap! [to perform repeated minimum computations]

- Use keys = frequencies
 - After extracting the two smallest-frequency symbols, re-Insert the new meta-symbol [new key = sum of the 2 old ones]
- ⇒ Iterative, $O(n \log n)$ implementation.

Even faster: (Non-trivial exercise) Sorting + $O(n)$ additional work.

- Manage (meta-)symbols using two queues.



Huffman Codes

Correctness Proof

Algorithms: Design
and Analysis, Part II


Correctness of Huffman's Algorithm

Theorem: [Huffman 52] Huffman's algorithm computes a binary tree (with leaves \leftrightarrow symbols of Σ) that minimizes the average encoding length

$$L(T) = \sum_{i \in \Sigma} p_i [\text{depth of leaf } i \text{ in } T].$$

Proof: By induction on $n = |\Sigma|$. (Can assume $n \geq 2$.)

Base case: When $n = 2$, algorithm outputs the optimal tree.
(Needs 1 bit per symbol)




Inductive step: Fix input with $n = |\Sigma| > 2$.

By inductive hypothesis: Algorithm solves smaller subproblems (for Σ') optimally.

Inductive Step

Let $\Sigma' = \Sigma$ with a, b (symbols with smallest frequencies) replaced by meta-symbol ab . Define $p_{ab} = p_a + p_b$.

Recall: Exact correspondence between:



Important: For every such pair T' and T , $L(T) - L(T')$ is (after cancellation)

$$p_a [a's \text{ depth in } T] + p_b [b's \text{ depth in } T] - p_{ab} [ab's \text{ depth in } T'] =$$

Each is one more than

$$= p_a(d+1) + p_b(d+1) - (p_a + p_b)d = p_a + p_b, \text{ Independent of } T, T'!$$

Proof of Theorem

Inductive hypothesis: Huffman's algorithm computes a tree \hat{T}' that minimizes $L(T')$ for Σ' .

Upshot of last slide: Corresponding tree \hat{T} minimizes $L(T)$ for Σ over all trees in X_{ab} (i.e., where a & b are siblings)


Key lemma: [Completes proof of theorem] There is an optimal tree (for Σ) in X_{ab} . [i.e., a & b were “safe” to merge]

Intuition: Can make an optimal tree better by pushing a & b as deep as possible (since a, b have smallest frequencies).

Proof of Key Lemma

By exchange argument. Let T^* be any tree that minimizes $L(T)$ for Σ . Let x, y be siblings at the deepest level of T^* .

The exchange: Obtain \hat{T} from T^* by swapping $a \leftrightarrow x$, $b \leftrightarrow y$



Note: $\hat{T} \in X_{ab}$ (by choice of x, y).

To finish: Will show that $L(\hat{T}) \leq L(T^*)$

$\Rightarrow \hat{T}$ also optimal, completes proof]

Reason:

$$\begin{aligned}
 L(T^*) - L(\hat{T}) &= (p_x - p_a) \text{ [x's depth in } T^* - a\text{'s depth in } T^*] \\
 &+ (p_y - p_b) \text{ [y's depth in } T^* - b\text{'s depth in } T^*] \\
 &\geq 0 \quad \text{QED!}
 \end{aligned}$$

≥ 0 since a, b have smallest frequencies

≥ 0 by choice of x, y

Notes on Running Time

Naive implementation: $O(n^2)$ time, where $n = |\Sigma|$.

Speed ups: - Use a heap! [to perform repeated minimum computations]

- Use keys = frequencies
 - After extracting the two smallest-frequency symbols, re-Insert the new meta-symbol [new key = sum of the 2 old ones]
- ⇒ Iterative, $O(n \log n)$ implementation.

Even faster: (Non-trivial exercise) Sorting + $O(n)$ additional work.

- Manage (meta-)symbols using two queues.