



All-Pairs Shortest Paths (APSP)

Algorithms: Design
and Analysis, Part II

Problem Definition

Problem Definition

Input: Directed graph $G = (V, E)$ with edge costs c_e for each edge $e \in E$, [No distinguished source vertex.]

Goal: Either

(A) Compute the length of a shortest $u \rightarrow v$ path for all pairs of vertices $u, v \in V$

OR

(B) Correctly report that G contains a negative cycle.

Quiz

Question: How many invocations of a single-source shortest-path subroutine are needed to solve the all-pairs shortest path problem?

[$n = \#$ of vertices]

A) 1

B) $n - 1$

C) n

D) n^2

Running time (nonnegative edge costs):

$$n \cdot \text{Dijkstra} = O(nm \log n) = \begin{cases} O(n^2 \log n) & \text{if } m = \Theta(n) \\ O(n^3 \log n) & \text{if } m = \Theta(n^2) \end{cases}$$

Running time (general edge costs):

$$n \cdot \text{Bellman-Ford} = O(n^2 m) = \begin{cases} O(n^3) & \text{if } m = \Theta(n) \\ O(n^4) & \text{if } m = \Theta(n^2) \end{cases}$$



All-Pairs Shortest Paths (APSP)

Algorithms: Design
and Analysis, Part II

Optimal Substructure

Motivation

Floyd-Warshall algorithm: $O(n^3)$ algorithm for APSP.

- Works even with graphs with negative edge lengths.

Thus: (1) At least as good as n Bellman-Fords, better in dense graphs.

(2) In graphs with nonnegative edge costs, competitive with n Dijkstra's in dense graphs.

Important special case: Transitive closure of a binary (i.e., all-pairs reachability) relation.

Open question: Solve APSP significantly faster than $O(n^3)$ in dense graphs?

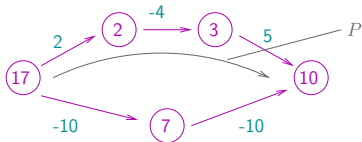
Optimal Substructure

Recall: Can be tricky to define ordering on subproblems in graph problems.

Key idea: Order the vertices $V = \{1, 2, \dots, n\}$ arbitrarily. Let $V^{(k)} = \{1, 2, \dots, k\}$.

Lemma: Suppose G has no negative cycle. Fix source $i \in V$, destination $j \in V$, and $k \in \{1, 2, \dots, n\}$. Let $P =$ shortest (cycle-free) i - j path with all internal nodes in $V^{(k)}$.

Example: $[i = 17, j = 10, k = 5]$



Optimal Substructure (con'd)


Optimal substructure lemma: Suppose G has no negative cost cycle. Let P be a shortest (cycle-free) i - j path with all internal nodes in $V^{(k)}$. Then:

Case 1: If k not internal to P , then P is a shortest (cycle-free) i - j path with all internal vertices in $V^{(k-1)}$.

Case 2: If k is internal to P , then:

P_1 = shortest (cycle-free) i - k path with all internal nodes in $V^{(k-1)}$ and

P_2 = shortest (cycle-free) k - j path with all internal nodes in $V^{(k-1)}$



Proof: Similar to Bellman-Ford opt substructure (you check!)



All-Pairs Shortest Paths (APSP)

Algorithms: Design
and Analysis, Part II

The Floyd-Warshall
Algorithm

Quiz

Setup: Let $A = 3\text{-D array (indexed by } i, j, k\text{)}.$

Intent: $A[i, j, k]$ = length of a shortest i - j path with all internal nodes in $\{1, 2, \dots, k\}$ (or $+\infty$ if no such paths)

Question: What is $A[i, j, 0]$ if

(1) $i = j$ (2) $(i, j) \in E$ (3) $i \neq j$ and $(i, j) \notin E$

A) 0, 0, and $+\infty$

B) 0, c_{ij} , and c_{ij}

C) 0, c_{ij} , and $+\infty$

D) $+\infty$, c_{ij} , and $+\infty$

The Floyd-Warshall Algorithm

Let A = 3-D array (indexed by i, j, k)

Base cases: For all $i, j \in V$:

$$A[i, j, 0] = \left\{ \begin{array}{ll} 0 & \text{if } i = j \\ c_{ij} & \text{if } (i, j) \in E \\ +\infty & \text{if } i \neq j \text{ and } (i, j) \notin E \end{array} \right\}$$

For $k = 1$ to n

For $i = 1$ to n

For $j = 1$ to n


$$A[i, j, k] = \min \left\{ \begin{array}{ll} A[i, j, k-1] & \text{Case 1} \\ A[i, k, k-1] + A[k, j, k-1] & \text{Case 2} \end{array} \right\}$$

Correctness: From optimal substructure + induction, as usual.

Running time: $O(1)$ per subproblem, $O(n^3)$ overall.

Odds and Ends

Question #1: What if input graph G has a negative cycle?



Answer: Will have $A[i, i, n] < 0$ for at least one $i \in V$ at end of algorithm.

Question #2: How to reconstruct a shortest i - j path?

Answer: In addition to A , have Floyd-Warshall compute $B[i, j] = \max$ label of an internal node on a shortest i - j path for all $i, j \in V$.

[Reset $B[i, j] = k$ if 2nd case of recurrence used to compute $A[i, j, k]$]

\Rightarrow Can use the $B[i, j]$'s to recursively reconstruct shortest paths!



All-Pairs Shortest Paths (APSP)

Algorithms: Design
and Analysis, Part II

A Reweighting
Technique

Motivation

Recall: APSP reduces to n invocations of SSSP.

- Nonnegative edge lengths: $O(mn \log n)$ via Dijkstra

- General edge lengths: $O(mn^2)$ via Bellman-Ford

Johnson's algorithm: Reduces APSP to

- 1 invocation of Bellman-Ford ($O(mn)$)

- n invocations of Dijkstra ($O(nm \log n)$)

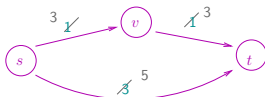
Running time: $O(mn) + O(mn \log n) = O(mn \log n)$

As good as with nonnegative edge lengths!

Quiz

Suppose: $G = (V, E)$ directed graph with edge lengths. Obtain G' from G by adding a constant M to every edge's length. When is the shortest path between a source s and a destination t guaranteed to be the same in G and G' ?

- A) When G has no negative-cost cycle
- B) When all edge costs of G are nonnegative
- C) When all s - t paths in G have the same number of edges
- D) Always

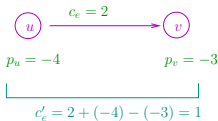


Quiz

Setup: $G = (V, E)$ is a directed graph with general edge lengths c_e . Fix a real number p_v for each vertex $v \in V$.

Definition: For every edge $e = (u, v)$ of G , $c'_e := c_e + p_u - p_v$

Question: If the s - t path P has length L with the original edge lengths $\{c_e\}$, what is P 's length with the new edge length $\{c'_e\}$?



- A) L
- B) $L + p_s + p_t$
- C) $L + p_s - p_t$
- D) $L - p_s + p_t$

$$\text{New length} = \sum_{e \in P} c'_e = \sum_{e=(u,v) \in P} [c_e + p_u - p_v] = (\sum_{e \in P} c_e) + p_s - p_t$$

Reweighting

Summary: Reweighting using vertex weights $\{p_v\}$ adds the same amount (namely, $p_s - p_t$) to every s - t path.

Consequence: Reweighting always leaves the shortest path unchanged.

Why useful? What if:

- (1) G has some negative edge lengths
- (2) After reweighting by some $\{p_v\}$, all edge lengths become nonnegative!

Question: Do such weights always exist?

Yes, and can be computed using the Bellman-Ford algorithm!

Requires Bellman-Ford, enables Dijkstra!

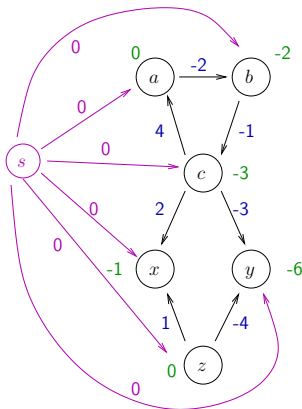


All-Pairs Shortest Paths (APSP)

Algorithms: Design
and Analysis, Part II

Johnson's Algorithm


Example



Note: Adding s does not add any new u - v paths for any $u, v \in G$.

Key insight: Define vertex weight $p_v :=$ length of a shortest s - v path.

Example (con'd)




Recall: For each edge $e = (u, v)$, define $c'_e = c_e + p_u - p_v$.

Note: After reweighting, all edge lengths nonnegative! \Rightarrow Can compute all (reweighted) shortest paths via n Dijkstra computations! [No need for Bellman-Ford]

Johnson's Algorithm

Input: Directed graph $G = (V, E)$, general edge lengths c_e .

- (1) Form G' by adding a new vertex s and a new edge (s, v) with length 0 for each $v \in G$.



- (2) Run Bellman-Ford on G' with source vertex s . [If B-F detects a negative-cost cycle in G' (which must lie in G), halt + report this.]
- (3) For each $v \in G$, define p_v = length of a shortest $s \rightarrow v$ path in G' . For each edge $e = (u, v) \in G$, define $c'_e = c_e + p_u - p_v$.
- (4) For each vertex u of G : Run Dijkstra's algorithm in G , with edge lengths $\{c'_e\}$, with source vertex u , to compute the shortest-path distance $d'(u, v)$ for each $v \in G$.
- (5) For each pair $u, v \in G$, return the shortest-path distance $d(u, v) := d'(u, v) - p_u + p_v$

Analysis of Johnson's Algorithm

Running time: $O(n) + O(mn) + O(m) + O(nm \log n) + O(n^2)$

Step (1), form G' Step (2), run BF Step (3), form c' Step (4), n Dijkstra Step (5), $O(1)$ work per $u-v$ pair

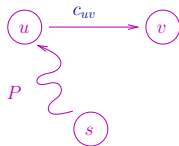
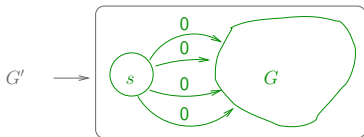
$= O(mn \log n)$. [Much better than Floyd-Warshall for sparse graphs!]

Correctness: Assuming $c'_e \geq 0$ for all edges e (see next slide for proof), correctness follows from last video's quiz.

[Reweighting doesn't change the shortest $u-v$ path, it just adds $(p_u - p_v)$ to its length]

Correctness of Johnson's Algorithm

Claim: For every edge $e = (u, v)$ of G , the reweighted length $c'_e = c_e + p_u - p_v$ is nonnegative.



Proof: Fix an edge (u, v) . By construction,
 p_u = length of a shortest s - u path in G'
 p_v = length of a shortest s - v path in G'
Let P = a shortest s - u path in G' (with length p_u - exists, by construction of G')

$\Rightarrow P + (u, v)$ = an s - v path with length $p_u + c_{uv}$

\Rightarrow Shortest s - v path only shorter, so $p_v \leq p_u + c_{uv}$

$\Rightarrow c'_{uv} = c_{uv} + p_u - p_v \geq 0$. QED!