



Algorithms: Design
and Analysis, Part II

Local Search

The Maximum Cut Problem

The Maximum Cut Problem

Input: An undirected graph $G = (V, E)$.

Goal: A cut (A, B) – a partition of V into two non-empty sets – that maximizes the number of crossing edges.


Sad fact: NP-complete.

Computationally tractable special case: Bipartite graphs (i.e., where there is a cut such that all edges are crossing)

Exercise: Solve in linear time via breadth-first search

Quiz


Question: What is the value of a maximum cut in the following graph?



- A) 4
- B) 6
- C) 8
- D) 10

A Local Search Algorithm

Notation: For a cut (A, B) and a vertex v , define
 $c_v(A, B) = \#$ of edges incident on v that cross (A, B)
 $d_v(A, B) = \#$ of edges incident on v that don't cross (A, B)



Local search algorithm:

- (1) Let (A, B) be an arbitrary cut of G .
- (2) While there is a vertex v with $d_v(A, B) > c_v(A, B)$:
 - Move v to other side of the cut


[key point: increases number of crossing edges by $d_v(A, B) - c_v(A, B) > 0$]
- (3) Return final cut (A, B)

Note: Terminates within $\binom{n}{2}$ iterations [+ hence in polynomial time].

Performance Guarantees

Theorem: This local search algorithm always outputs a cut in which the number of crossing edges is at least 50% of the maximum possible. (Even 50% of $|E|$)

Tight example:



Cautionary point: Expected number of crossing edges of a random cut already is $\frac{1}{2}|E|$.

Proof: Consider a random cut (A, B) . For edge $e \in E$, define

$$X_e = \begin{cases} 1 & \text{if } e \text{ crosses } (A, B) \\ 0 & \text{otherwise} \end{cases}.$$

We have $E[X_e] = \Pr[X_e = 1] = 1/2$.
So $E[\# \text{ crossing edges}] = E[\sum_e X_e] = \sum_e E[X_e] = |E|/2$. QED

Proof of Performance Guarantee

Let (A, B) be a locally optimal cut. Then, for every vertex v , $d_v(A, B) \leq c_v(A, B)$. Summing over all $v \in V$:

$$\sum_{v \in V} d_v(A, B) \leq \sum_{v \in V} c_v(A, B)$$

counts each non-crossing edge twice counts each crossing edge twice

So:

$$2 \cdot [\# \text{ of non-crossing edges}] \leq 2 \cdot [\# \text{ of crossing edges}]$$

$$2 \cdot |E| \leq 4 \cdot [\# \text{ of crossing edges}]$$

$$\# \text{ of crossing edges} \geq \frac{1}{2}|E| \quad \text{QED!}$$

The Weighted Maximum Cut Problem

Generalization: Each edge $e \in E$ has a nonnegative weight w_e , want to maximize total weight of crossing edges.

Notes:

- (1) Local search still well defined
- (2) Performance guarantee of 50% still holds for locally optimal cuts [you check!] (also for a random cut)
- (3) No longer guaranteed to converge in polynomial time [non-trivial exercise]



Algorithms: Design
and Analysis, Part II

Local Search

Principles of Local Search

Neighborhoods

Let X = set of candidate solutions to a problem.

Examples: Cuts of a graph, TSP tours, CSP variable assignments


Key ingredient: Neighborhoods

- For each $x \in X$, specify which $y \in X$ are its “neighbors”

Examples: x, y are neighboring cuts \iff Differ by moving one vertex

x, y are neighboring variable assignments \iff Differ in the value of a single variable

x, y are neighboring TSP tours \iff Differ by 2 edges



A Generic Local Search Algorithm

- (1) Let x = some initial solution.
- (2) While the current solution x has a superior neighboring solution y :
 Set $x := y$
- (3) Return the final (locally optimal) solution x

FAQ

Question: How to pick initial solution x ?

Answer #1: Use a random solution.

⇒ Run many independent trials of local search, return the best locally optimal solution found.

Answer #2: Use your best heuristics
(i.e., use local search as a postprocessing step to make your solution even better).

Question #2: If there are superior neighboring y , which to choose?

Possible answers: (1) Choose at random, (2) biggest improvement, (3) more complex heuristics.

Question #3: How to define neighborhoods?

Note bigger neighborhoods ⇒ slower to verify local optimality, but fewer (bad) local optima

Answer: Find “sweet spot” between solution quality and efficient searchability.

FAQ II

Question: Is local search guaranteed to terminate (eventually)?

Answer: If X is finite and every local step improves some objective function, then yes.

Question: Is local search guaranteed to converge quickly?

Answer: Usually not. [though it often does in practice] (see “smoothed analysis”)

Question: Are locally optimal solutions generally good approximations to globally optimal ones?

Answer: No. [To mitigate, run randomized local search many times, remember the best locally optimal solution found]



Local Search

The 2-SAT Problem

Algorithms: Design
and Analysis, Part II

2-SAT

Input:

- (1) n Boolean variables x_1, x_2, \dots, x_n . (Can be set to TRUE or FALSE)
- (2) m clauses of 2 literals each (“literal” = x_i or $\neg x_i$)

Example: $(x_1 \vee x_2) \wedge (\neg x_1 \vee x_3) \wedge (x_3 \vee x_4) \wedge (\neg x_2 \vee \neg x_4)$

Output: “Yes” if there is an assignment that simultaneously satisfies every clause, “no” otherwise.

Example: “yes”, via (e.g.) $x_1 = x_3 = \text{TRUE}$ and $x_2 = x_4 = \text{FALSE}$

(In)Tractability of SAT

2-SAT: Can be solved in polynomial time!

- Reduction to computing strongly connected components (nontrivial exercise)
- “Backtracking” works in polynomial time (nontrivial exercise)
- Randomized local search (next)

3-SAT: Canonical NP-complete

- Brute-force search $\approx 2^n$ time
- Can get time $\approx \left(\frac{4}{3}\right)^n$ via randomized local search [Schöningh '02]

Papadimitriou's 2-SAT Algorithm

Repeat $\log_2 n$ times:

- Choose random initial assignment
- Repeat $2n^2$ times:
 - If current assignment satisfies all clauses, halt + report this
 - Else, pick arbitrary unsatisfied clause and flip the value of one of its variables [choose between the two uniformly at random]

Report “unsatisfiable”

Key question: If there's a satisfying assignment, will the algorithm find one (with probability close to 1)?

Obvious good points:

- (1) Runs in polynomial time
- (2) Always correct on unsatisfiable instances



Algorithms: Design
and Analysis, Part II

Local Search


Random Walks on a
Line

Random Walks

Key to analyzing Papadimitriou's algorithm:

Random walks on the nonnegative integers (trust me!)

Setup: Initially (at time 0), at position 0.



At each time step, your position goes up or down by 1, with 50/50 probability.

[Except if at position 0, in which case you move to position 1 with 100% probability]

Quiz

Notation: For an integer $n \geq 0$, let T_n = number of steps until random walk reaches position n .

[A random variable, sample space = coin flips at all time steps]

Question: What is $E[T_n]$? (your best guess)

A) $\Theta(n)$

B) $\Theta(n^2)$

C) $\Theta(n^3)$

D) $\Theta(2^n)$

Coming up: $E[T_n] = n^2$.

Analysis of T_n

Let Z_i = number of random walk steps to get to n from i . (Note $Z_0 = T_n$)

Edge cases: $E[Z_n]=0$, $E[Z_0]=1+E[Z_1]$

For $i \in \{1, 2, \dots, n-1\}$

$$\begin{aligned} E[Z_i] &= \Pr[\text{go left}] E[Z_i \mid \text{go left}] + \Pr[\text{go right}] E[Z_i \mid \text{go right}] \\ &= 1 + \frac{1}{2} E[Z_{i+1}] + \frac{1}{2} E[Z_{i-1}] \end{aligned}$$

Diagram annotations: Arrows point from the text above to terms in the equation. From '1/2' to 'Pr[go left]'. From '(1+E[Z_{i-1}])' to 'E[Z_i | go left]'. From '1/2' to 'Pr[go right]'. From '(1+E[Z_{i+1}])' to 'E[Z_i | go right]'.

Rearranging: $E[Z_i] - E[Z_{i+1}] = E[Z_{i-1}] - E[Z_i] + 2$

Finishing the Proof of Claim

So:

$$\begin{array}{rcl} E[Z_0] - E[Z_1] & = & 1 \\ E[Z_1] - E[Z_2] & = & 3 \\ E[Z_2] - E[Z_3] & = & 5 \\ & \vdots & \\ + E[Z_{n-1}] - E[Z_n] & = & 2n - 1 \end{array}$$

Diagram: Purple curved arrows group the terms on the right side of the equations. The first arrow groups the 1 and 3, the second groups the 3 and 5, and the third groups the 5 and 7. A text label points to these groups.

$\frac{n}{2}$ pairs of numbers,
each sums to $2n$

$$E[Z_0] = n^2$$

||

$$E[T_n]$$

QED!

A Corollary

Corollary: $\Pr[T_n > 2n^2] \leq \frac{1}{2}$. (Special case of Markov's inequality)

Proof: Let p denote $\Pr[T_n > 2n^2]$.

We have $n^2 = E[T_n]$

by last claim

$$= \sum_{k=0}^{2n^2} k \Pr[T_n = k] + \sum_{k=2n^2+1}^{\infty} k \Pr[T_n = k]$$

$$\geq 2n^2 \Pr[T_n > 2n^2]$$

$$= 2n^2 p.$$

$$\Rightarrow p \leq \frac{1}{2} \quad \text{QED!}$$



Algorithms: Design
and Analysis, Part II

Local Search

Analysis of
Papadimitriou's Algorithm

Papadimitriou's Algorithm

n = number of variables

Repeat $\log_2 n$ times:

- Choose random initial assignment
- Repeat $2n^2$ times:
 - If current assignment satisfies all clauses, halt + report this
 - Else, pick arbitrary unsatisfied clause and flip the value of one of its variables [choose between the two uniformly at random]

Report “unsatisfiable”

Obvious good points:

- (1) Runs in polynomial time
- (2) Always correct on unsatisfiable instances

Satisfiable Instances

Theorem: For a satisfiable 2-SAT instance with n variables, Papadimitriou's algorithm produces a satisfying assignment with probability $\geq 1 - \frac{1}{n}$.

Proof: First focus on a single iteration of the outer for loop.

Fix an arbitrary satisfying assignment a^* .

Let a_t = algorithm's assignment after inner iteration t
($t = 0, 1, \dots, 2n^2$) [a random variable]

Let X_t = number of variables on which a_t and a^* agree.
($X_t \in \{0, 1, \dots, n\}$)

Note: If $X_t = n$, algorithm halts with satisfying assignment a^* .

Proof of Theorem (con'd)

Key point: Suppose a_t not a satisfying assignment and algorithm picks unsatisfied clause with variables x_i, x_j .

Note: Since a^* is satisfying, it makes a different assignment than x_i or x_j (or both).

Consequence of algorithm's random variable flip:

(1) If a^* and a_t differ on both x_i & x_j , then $X_{t+1} = X_t + 1$ (100% probability)

(2) If a^* and a_t differ on exactly one of x_i, x_j , then

$$X_{t+1} = \begin{cases} X_t + 1 & (50\% \text{ probability}) \\ X_t - 1 & (50\% \text{ probability}) \end{cases}$$

Quiz: Connection to Random Walks

Question: The random variables $X_0, X_1, \dots, X_{2n^2}$ behave just like a random walk of the nonnegative integers except that:



- A) Sometimes move right with 100% probability (instead of 50%)
- B) Might have $X_0 > 0$ instead of $X_0 = 0$
- C) Might stop early, before $X_t = n$
- D) All of the above

Completing the Proof

Consequence: Probability that a single iteration of the outer for loop finds a satisfying assignment is $\geq \Pr[T_n \leq 2n^2] \geq 1/2$

from last video



Thus:

$$\begin{aligned}\Pr[\text{algorithm fails}] &\leq \Pr[\text{all } \log_2 n \text{ independent trials fail}] \\ &\leq \left(\frac{1}{2}\right)^{\log_2 n} \\ &= \frac{1}{n}. \quad \text{QED!}\end{aligned}$$